

## (iv) non-atomic state transitions

- In the protocol state diagram, all transitions are assumed to take place instantaneously/atomically
- But each change involves several intermediate actions, even if the bus is atomic
- Actions like:
  - Cache look-up, bus arbitration, actions taken by other controllers at their caches, action of issuing processor's controller, final block write
- Can have race conditions among components of different operations
- Ex: Request for block 'B' may come to some processor, while the processor is still changing state for block 'B'

# Issues: non-atomic ...

Suppose P1 and P2 attempt to write cached block 'A' simultaneously

- (1) P1 has blk-A in cache
- (2) P1 wants to write, therefore sends BusRdX or BusUpgr and then S--> M
- (4) P1 sees P2's request for blk-A. Cancels own request makes state S-->I
- (5) P1 will need to send new BusRdX to get fresh copy from P2

- (1) P2 does not have blk-A
- (2) P2 wants to write, therefore sends BusRdX
- (3) P2 wins Bus arbitration
- (5) P2 modifies blk-A



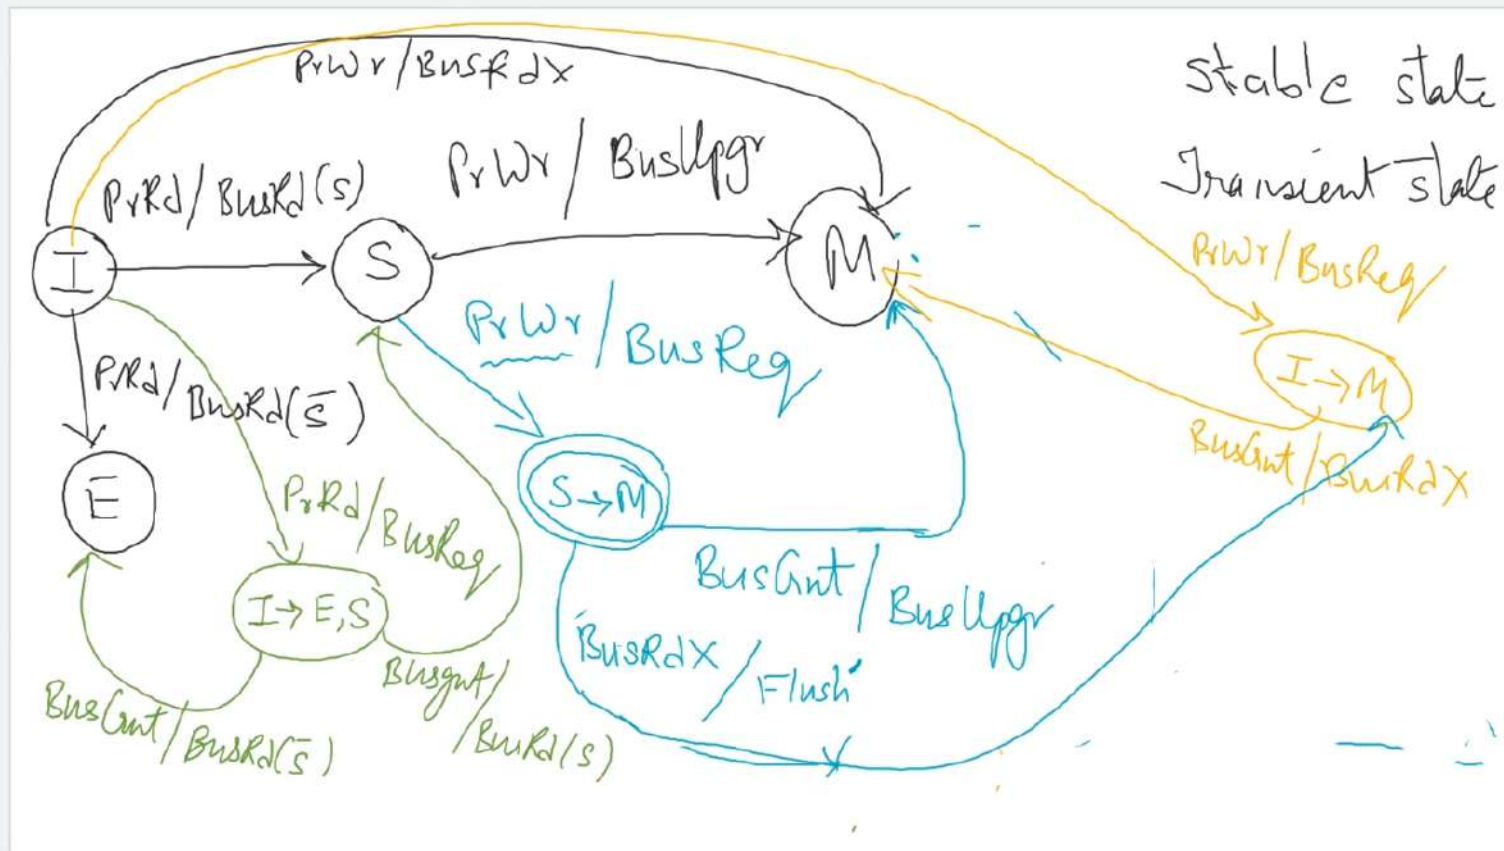
23\_Aug\_2021

↪



Set background

Clear frame



# Transient states

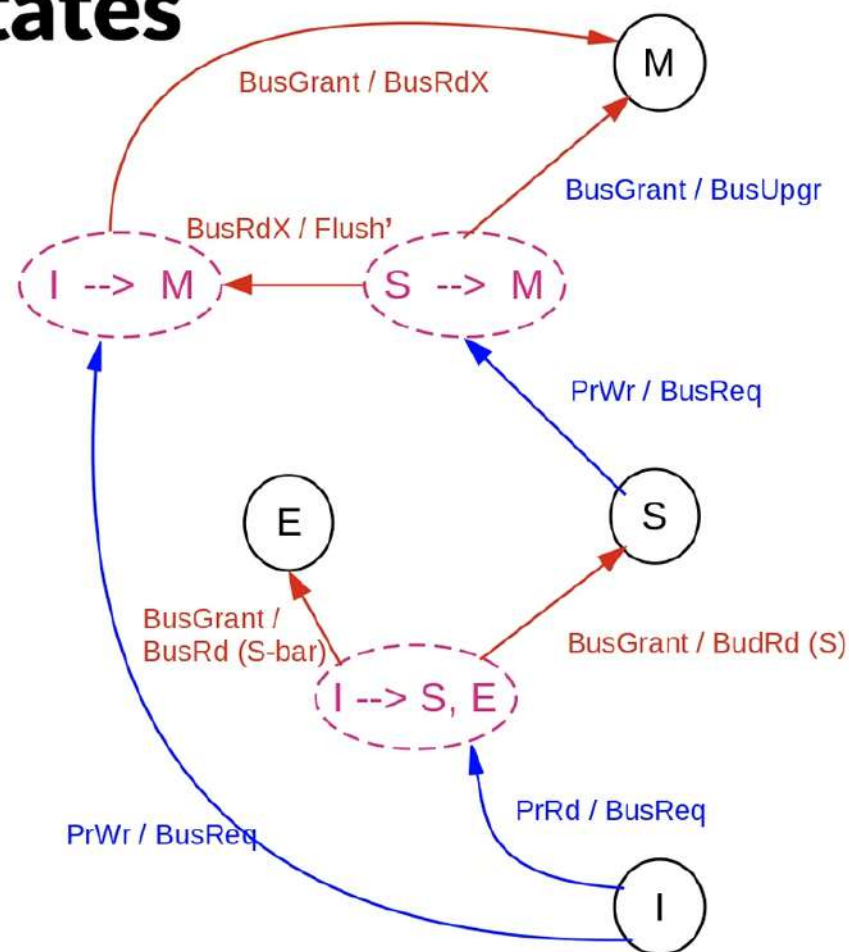
- Changes in MESI:

$I \xrightarrow{\text{red}} S, E$

$S \xrightarrow{\text{blue}} M$

$I \xrightarrow{\text{blue}} M$

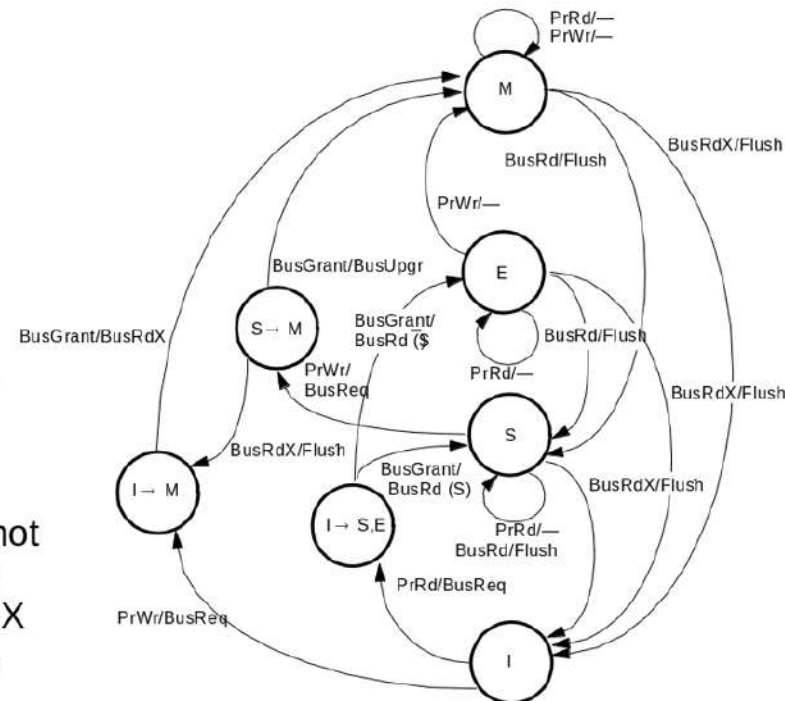
- Shown dotted in the FSM





# Handling Non-atomicity: Transient States

- 2 types of states
  - Stable (e.g. original MESI)
  - Transient or Intermediate
- Increases complexity
  - e.g. don't send BusUpgr, rather use other mechanisms to avoid data transfer
- If we always use BusRdX then we do not need the intermediate state (S→M), as the output action will always be BusRdX
  - In case of BusUpgr we have to modify the output action to BusRdX and hence one transient state is required.



# Serialisation

- Processor-cache handshake must preserve the order determined by the serialisation of bus transactions
- On a write to 'S' block
  - Tempting to start writing before acquiring 'M' ownership .... processor need not wait
  - Because there will be problems, as other transactions (conflicting) may appear in-between
- If other transaction completed before this one, states may be inconsistent
- These transactions must appear to the processor as occurring before the write (by this processor) since that is how they are serialised by bus and appear to other processors




# Write Completion for SC

- Need not wait for invalidations to actually happen. Just wait till it gets the bus
- Commit vs Complete
  - Don't know when invalidation actually inserted in destination process's local order
  - Only that it is before the next bus transaction and in same order for all processors
  - Local write-hits become visible to others only after another bus-transaction
  - What matters is not when the written data appears on the bus (wr-back), but when subsequent reads are guaranteed to see it
- Write atomicity
  - If read returns a value of write-W1 it means write transaction for W1 has already gone to the bus and therefore is completed (if it needed to)
- Deadlock



# Write Completion for SC

- Deadlock
  - Request-response protocols can lead to protocol-level fetch-deadlock
  - Therefore when attempting to issue requests, must service incoming transactions
  - Cache controller awaiting bus grant must snoop and even flush blocks
  - Else may not respond to requests that will release bus



P1 --> readEx on A to P2 . P1 holds B


P2 --> readEx on B to P1 . P2 holds A

Therefore P1, must keep readEx-A pending and serve 'B' to P2  
when P2 wins bus arbitration



# Write Completion for SC

- Deadlock
  - Request-response protocols can lead to protocol-level fetch-deadlock
  - Therefore when attempting to issue requests, must service incoming transactions
  - Cache controller awaiting bus grant must snoop and even flush blocks
  - Else may not respond to requests that will release bus



P1 --> readEx on A to P2 . P1 holds B

P2 --> readEx on B to P1 . P2 holds A

Therefore P1, must keep readEx-A pending and serve 'B' to P2  
when P2 wins bus arbitration

# Livelock and Starvation

- Many processors try to write same line/block
- Each one:
  - Obtains exclusive ownership via bus transaction (Assume block not in cache)
  - Block brought in cache and therefore processor tries to write it
  - If processor cache handshake is not properly designed...following is possible
    - At same time another processor requests same block ==> so I release it
- Livelock: I obtain ownership, but you steal it before I can write ...
  - Solution: don't let exclusive ownership be taken away before write is done
- Starvation: Fair arbitration on bus, FIFO service queues, FCFS service policies. Implementation needs buffers or counters to keep track how many times a processor served etc.



# Multi-level cache + Atomic Bus

/

Hemangee K. Kapoor

30

+23



MN

TU



YOGESH KUMAR

YK



DEVANSHI GUPTA

DG



NISHTHA SHARMA

NS



Syam Sanker

SS



TANVISH



Hemangee Kalpesh Kapoor



ASWATHY N S

AS

# Multi-level cache + Atomic Bus

Hemangee K. Kapoor

30

+23



MN

TU



YOGESH KUMAR

YK



DEVANSHI GUPTA

DG



NISHTHA SHARMA

NS



Syam Sanker

SS



TANVISH



Hemangee Kalpesh Kapoor

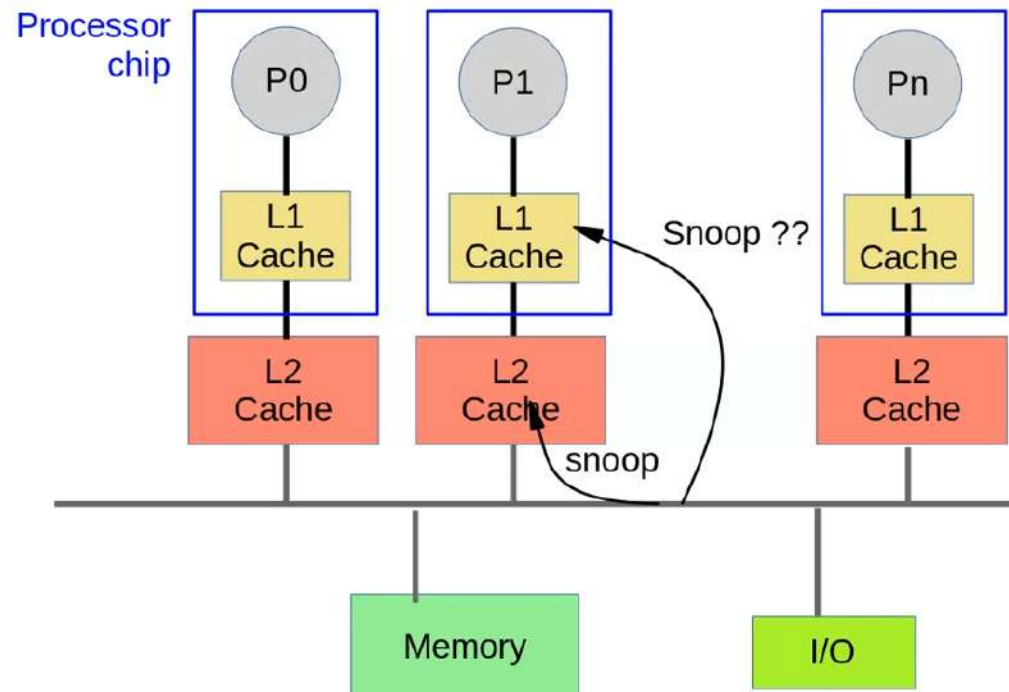


ASWATHY N S

AS



# Multi-Level Cache Hierarchies



# Multilevel cache hierarchies

- To handle multi-level cache one way is to have independent bus snooping hardware for each level of cache
- 3 reasons why this is ill-suited
  - L1 on processor-chip and snoop hardware will need extra pins to monitor the bus
  - Duplicating tags (of L1) will consume more area
  - Duplication of effort as most of time blocks in L1 are present in L2. Therefore snoop-by L1 un-necessary




# Multilevel cache hierarchies

- To handle multi-level cache one way is to have independent bus snooping hardware for each level of cache
- 3 reasons why this is ill-suited
  - L1 on processor-chip and snoop hardware will need extra pins to monitor the bus
  - Duplicating tags (of L1) will consume more area
  - Duplication of effort as most of time blocks in L1 are present in L2. Therefore snoop-by L1 un-necessary
- Solution: **snoop only at L2** and propagate relevant transactions. Needs inclusion property
- **Inclusion property:**
  - (1) L1 subset L2
  - (2) Any block in modified state in L1 is in marked-modified state in L2
  - 1 => All transactions relevant to L1 are relevant to L2
  - 2 => On BusRd, L2 can waive-off memory access and inform L1 to get updated copy of the block

# Maintaining Inclusion

- Maintaining inclusion is not trivial as L1 and L2 change state due to processor and bus respectively
- Two caches (L1 & L2) may choose to replace different block
  - (1) Set associative L1-cache with history-based replacement (ex LRU)
  - (2) Multiple caches at a level (ex: I\$ and D\$)
  - (3) Different cache block sizes
- But a common case automatically maintains inclusion (direct mapped)





# Set-assoc L1-cache with LRU


- L1 uses LRU --> 2-way set assoc, Blocks  $m_1, m_2, m_3$  fall in same set of L1  
L1 and L2 have same block size, L2 also 2-way set-assoc  
L2 is  $k$ -times larger than L1:  $n_2 = k * n_1$   
L1: set-0 :  $\langle m_1, m_2 \rangle$  ; L2: set- $i$  :  $\langle m_1, m_2 \rangle$

Suppose proc-P --> Read- $m_3$  --> **L1-miss** --> decides to **replace- $m_1$**  using LRU of L1

- as access pattern of L1 is decided by processor and L2 is un-aware of L1's LRU

**L2 decides to replace- $m_2$**


- If L2 was direct map -->  $m_1, m_2$  may fall in same set !
- **Inclusion can be violated if L1 is not-direct mapped and uses LRU replacement policy, no matter what is  $n_2, a_2, b_2$  in L2**



# Multiple caches at a level

- L1: Instr is in m1, Data is in m2
- Blocks may collide in L2
- L1:  $\langle I\$ : m1 \rangle$ ,  $\langle D\$ : m2 \rangle$
- L2: Direct map and m1, m2 map to same location
- L1: read-m1 ; L2: Load-m2
- L1: read-m2 ; L2: removes m1 and Load-m2. Due to this L2 does not have m1 but L1 has. Therefore violation of inclusion
- What if L2 is highly-associative cache?







L1

1 byte

	0
	1
	2
	3

set 0 0, 4  
set 1 1, 5  
set 2 2, 6  
set 3 3, 7  
Addr

L2

2B

	0
	1
	2
	3
	4
	5
	6
	7

set 0 = 0, 1  
set 1 = 2, 3  
set 7 =

# Ex: of violation - diff blk sizes

- $a_1, a_2=1$ ,  $b_1=1$ ,  $b_2=2$ ,  $n_1=4$ ,  $n_2=8$
- $L_1=4$  words,  $L_2=16$  words

--- L1 ---

Word locations

0,4,8 --> set-0

1,5,9 --> set-1

2, --> set-2

3, --> set-3

...

16 --> set-0

17 --> set-1

18, --> set-2

--- L2 ---

Word locations

0,1 --> set-0

2,3 --> set-1

4,5 --> set-2

6,7 --> set-3

...

14,15 --> set-7

16,17 --> set-0

18,19 --> set-1

- $L_1$  can have 0,17 at same time
- $L_2$  cannot ! As they map to same set

# Automatic inclusion

- L1 direct map:  $a1=1$
- L2 anything:  $a2 \geq 1$ , any replacement policy
- Same block size:  $b1=b2$
- New block put in both L1 and L2
- #sets in L1  $\leq$  #sets in L2:  $n1 \leq n2$
- This is a popular configuration to get the around problem of maintaining inclusion  
=> use L1 as direct map cache



# Preserving Inclusion Explicitly

- L2 takes coherence action of snooping on the bus
  - L1  $\setminus$  subset L2, Therefore certain transaction of L2 are relevant to L1
  - If L2 replaces block, information sent to L1 so that it can also remove those blocks
- Propagate bus transactions from L2 to L1
- Propagate modified state from L1 to L2 on writes?





# Preserving Inclusion Explicitly

- L2 takes coherence action of snooping on the bus
  - L1 subset L2, Therefore certain transaction of L2 are relevant to L1
  - If L2 replaces block, information sent to L1 so that it can also remove those blocks
- Propagate bus transactions from L2 to L1
  - Propagate all transactions?
    - Let L1 ignore if the mesg is irrelevant, Large number of such messages
    - Tags in L1 compared each time and so tags not available to CPU => bad idea
  - Use inclusion-bit?
    - L2 keeps track of which blocks are in L1 using inclusion-bit
    - Now L2 can selectively send requests to L1
- Propagate modified state from L1 to L2 on writes?
  - if L1 is write-through, just invalidate
    - L2 has updated copy ; Send the copy and invalidate block in L1
  - if L1 is write-back
    - add extra state-bit to L2 (dirty-but-stale); Get block from L1 before flush on the Bus

# Remarks

- Exclusion = no block is in more than any one cache
- Observation: if L2-design has restrictions implying difficult to maintain inclusion  
=> many new designs don't maintain inclusion  
Therefore snoop L1s

