

Split transaction bus + single level cache

Hemangee K. Kapoor

5

+19



AS



KARTIKEYA SAXENA



SHIVANGI KUMAR

SK

Syam Sankar

SS

SARASWATULA PHANI SAI PRANAV

SP

VARHADE AMEY ANANT

VA

Hemangee Kalpesh Kapoor



Hemangee Kalpesh Kapoor




Request-Response pairs

- BusRd = request ; data = response
- BusUpgr = request ; response = none
 - But it needs acknowledgement indicating that the transaction is committed and hence been serialised
 - To ensure that ACK does not appear on the bus as a separate transaction; the ACK is sent to the requesting processor by its own bus controller when the bus is granted to the BusUpgr request
- BusRdX = request ; response = data + ack of commitment (of inv by others)
 - The ack by others and data can be combined into one response transaction
- Write back = request ; response = none



Atomic vs Split

Atomic Transaction Bus



Split Transaction Bus




Atomic vs Split

- Design space too large
- When and how snoop results and data response provided
 - Are snoop results part of request phase or response phase?
- How many out-standing requests to allow
 - More requests => more buffering => complicated flow control
 - More requests => better bandwidth utilisation
- Are data response sent in the order of requests?
 - In-order: Intel Pentium Pro, DEC TurboLaser
 - Out-of-order: SGI Challenge, Sun Enterprise
- We will do a detailed split-bus example based on SGI Challenge Design

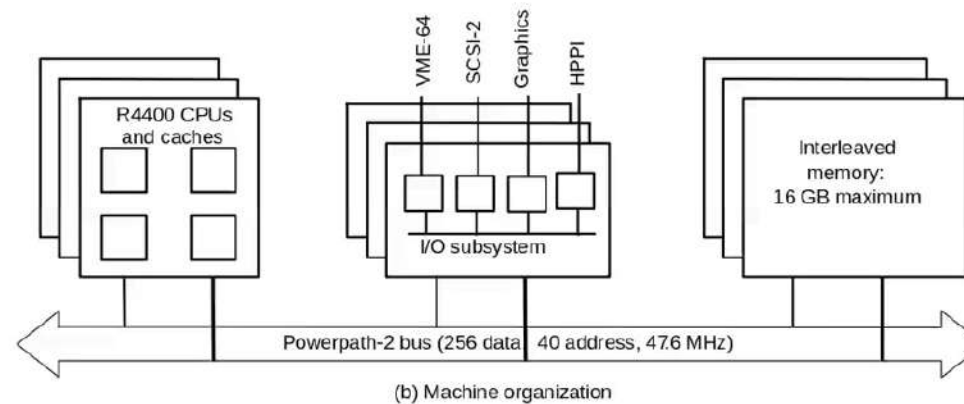


SGI Challenge overview

- 36 MIPS R4400 (peak 2.7 GFLOPS, 4 per board) or 18 MIPS R8000 (peak 5.4 GFLOPS, 2 per board)
- 8-way interleaved memory (up to 16 GB)
- 4 I/O busses of 320 MB/s each
- 1.2 GB/s Powerpath-2 bus @ 47.6 MHz, 16 slots, 329 signals
- 128 Bytes lines (1 + 4 cycles)
- Split-transaction with up to 8 outstanding reads
- all transactions take five cycles
- OS is variant of SVR4 UNIX called IRIX

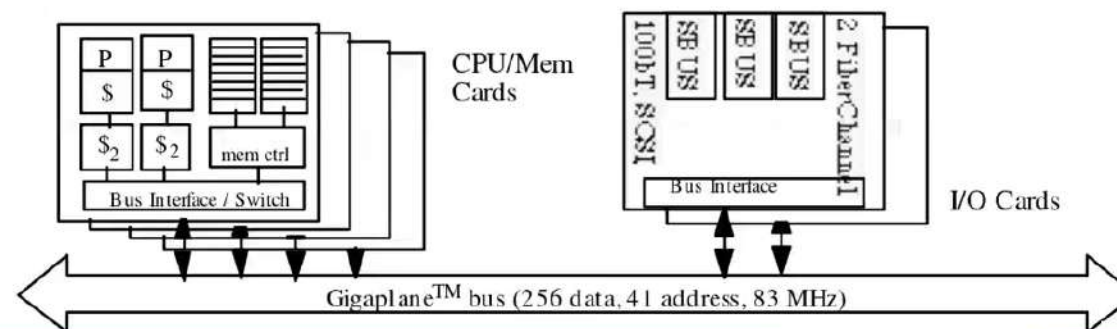


(a) A four-processor board



SUN Enterprise overview

- Up to 30 UltraSPARC processors (peak 9 GFLOPs)
- Gigaplane™ bus has peak bw 2.67 GB/s; upto 30GB memory
- 16 bus slots, for processing or I/O boards
- 2 CPUs and 1GB memory per board
- memory distributed, unlike Challenge, but protocol treats as centralized
- Each I/O board has 2 64-bit 25Mhz SBUSes
- OS is Solaris UNIX




Hemangee K. Kapoor

12

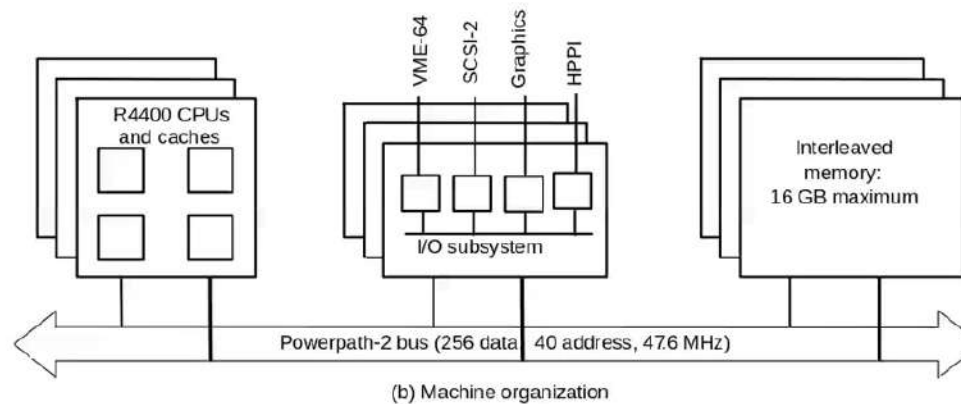


SGI Challenge overview

- 36 MIPS R4400 (peak 2.7 GFLOPS, 4 per board) or 18 MIPS R8000 (peak 5.4 GFLOPS, 2 per board)
- 8-way interleaved memory (up to 16 GB)
- 4 I/O busses of 320 MB/s each
- 1.2 GB/s Powerpath-2 bus @ 47.6 MHz, 16 slots, 329 signals
- 128 Bytes lines (1 + 4 cycles)
- Split-transaction with up to 8 outstanding reads
- all transactions take five cycles
- OS is variant of SVR4 UNIX called IRIX



(a) A four-processor board



SGI Challenge design

- No conflicting requests for the same block allowed on bus
 - At most 8 outstanding requests total, makes conflict detection tractable
- Flow control through negative acknowledgement (NACK)
 - If buffer full, NACK request so that requestor retries
 - Need separate lines for: command (inc. NACK), addr+tag, Data
- Responses may be in different order than requests
 - Total order among transactions maintained by bus requests (phase)
 - Snoop results go as part of response phase
- Look at ==>

Bus
Design

Snoop Results +
Conflicting requests


Flow
Control

Path of Request
Through system



Bus design and Req-Response matching

- Two separate buses, arbitrated independently
- (1) Request bus: command and address
- (2) Response bus: data
- Req-response are out-of-order
- Therefore need a tag to match 8 outstanding requests => 3-bit tag on 3-bit wide lines
 - Tag lines used during response phase
 - Therefore address lines are free for other new requests
- Separate bus line for arbitration and snoop result



Data response phase

- Data needs 5 cycles to transfer in our design
- Cache block = 128 bytes = 1024 bits
- Data bus = 256-bit wide
- For 1024 bits = $256 * 4$ cycles on data bus
- One additional cycle is for **turn-around time**
- Therefore $4+1 = 5$ cycles for data response phase
- The signalling technology requires this turn-around time between controllers driving the lines



3-Phases

(1) Address request: 5 cycles:

- Arbitration (for bus), Resolution (one wins, tag assigned), Address (send addr), Decode (cache-tag look-up, snoop-hit-check, make transaction visible to processor: state-change-E-S,M-S..etc), Ack (continue tag-lookup if not complete and extend ACK-cycle until snoop results are ready)
- At decode decide who (which cache) will respond (if any)
- Actual response will come later
- End of this it is known who will provide data: cache or memory

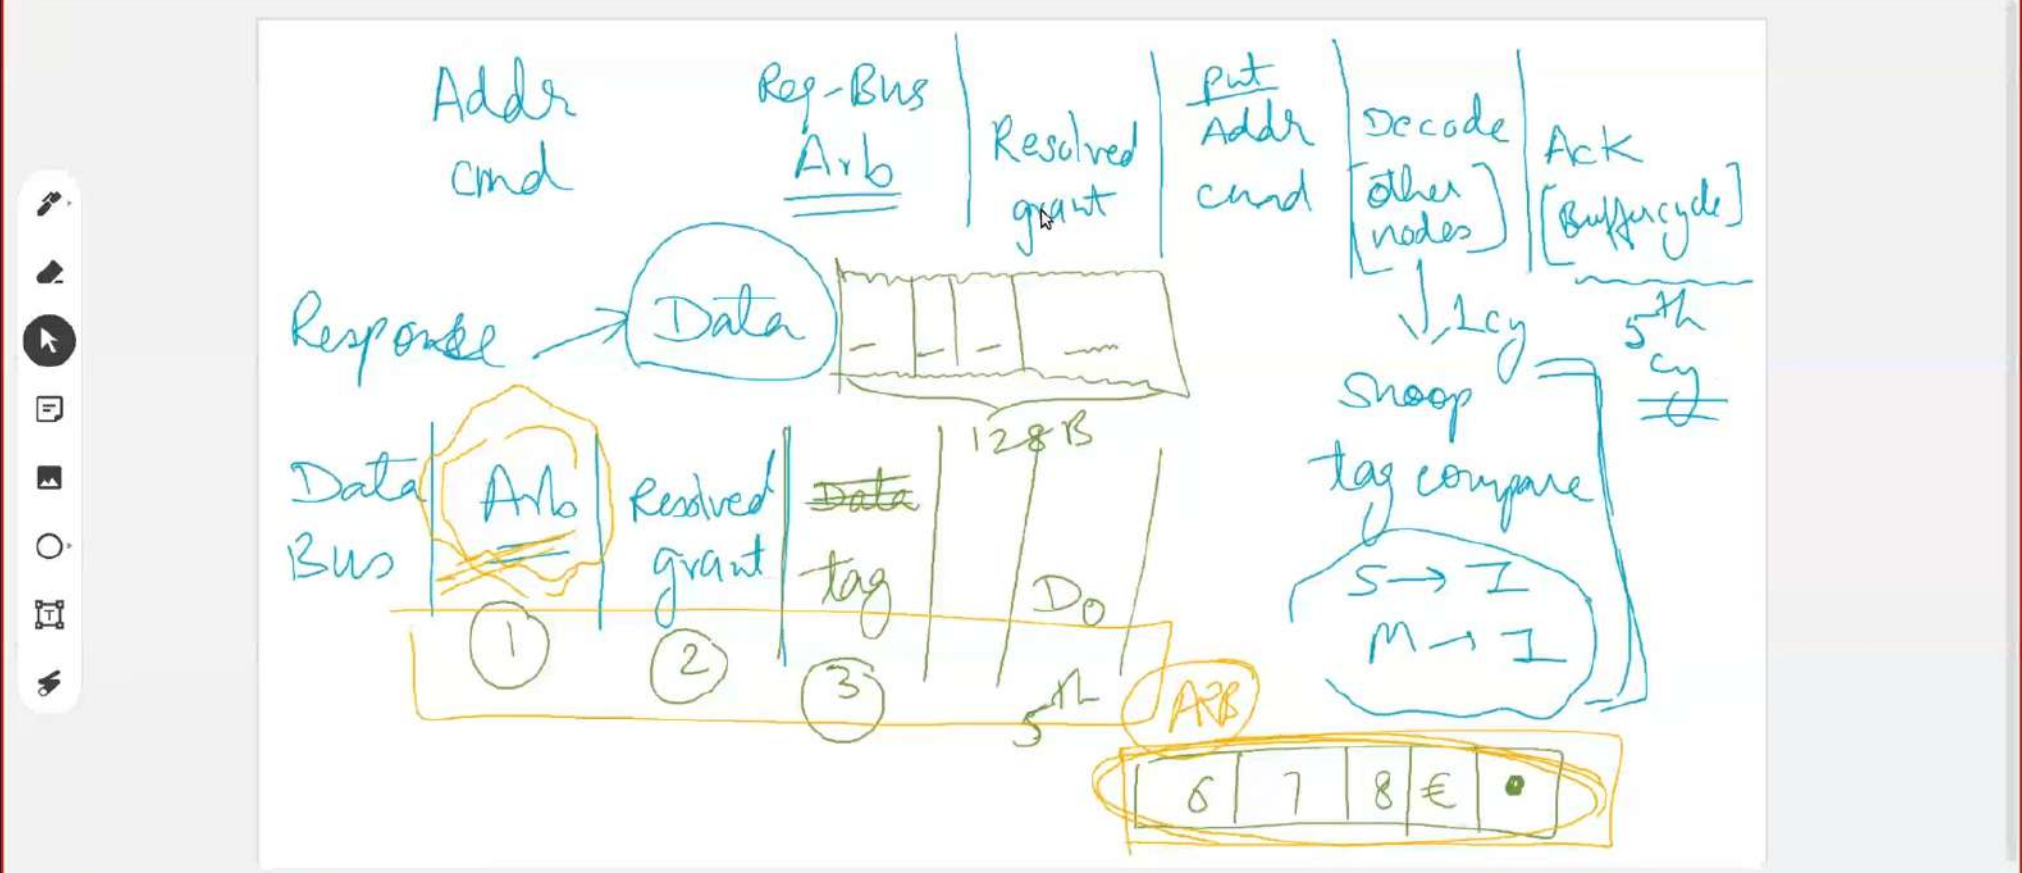
(2) Data request : 5 cycles

- Arbitration, Resolution, Address-cycle(=Tag-check), empty, [if target is ready then in ack cycle) Data-starts <D0,>

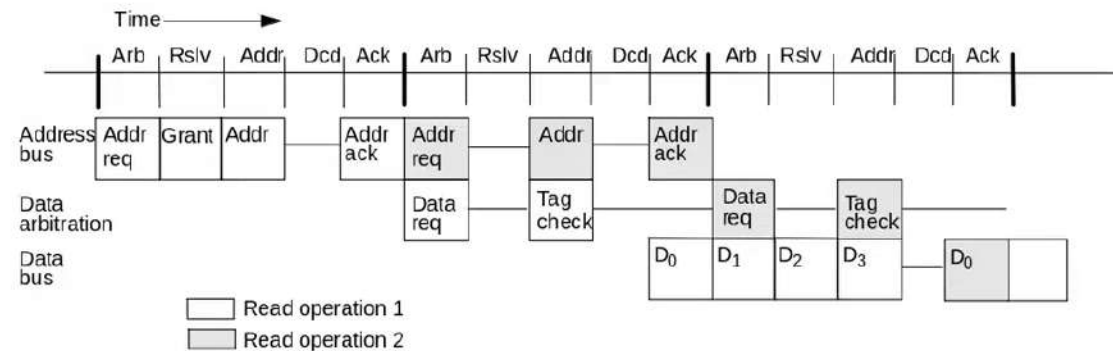
(3) Data response: 5 cycles (+ In parallel snoop results are conveyed so that states are correctly updated)

- Starts from ACK-stage of data -request phase
- Therefore D0 is sent in previous phase
- This phase sends <D1, D2, D3, empty,> <D0 of next request>
- The “empty” is for turn-around time





Ex: Read request that results in data transfer



- Scenario: cache has block in modified state
 - BusRd: send the block and state change E--> S
 - BusRdx or BusUpgr: send block and E --> I (inv.)



Example ...

P1 has block 'B'

P2 sends BusRd-req

P2: arbitrates [Addr-req, Grant, Addr.] <cy-1,2,3>] for addr-bus and sends addr of 'B' and command='BusRd' <say tag=1>

P1: sees the addr and does cache-tag compare (Decode <cy-4>) and send ACK (<cy-5>)

At cy-6 another processor pair needs to communicate

Say P3 <cy-6,7,8> wants block from P4 <cy-9,10>

Response phase of P1 <--> P2:

P1 will send block to P2

After request phase it is known which module will supply data: memory or cache



Example ..

P1 gets data-bus: cy-6,7

P1 sends the tag: <cy-8>

If target is ready then data block sent from <cy-10>

Data block = cy-10,11,12,13

Turn-around time = cy-14

Read operation-2 can send its data from cy-15

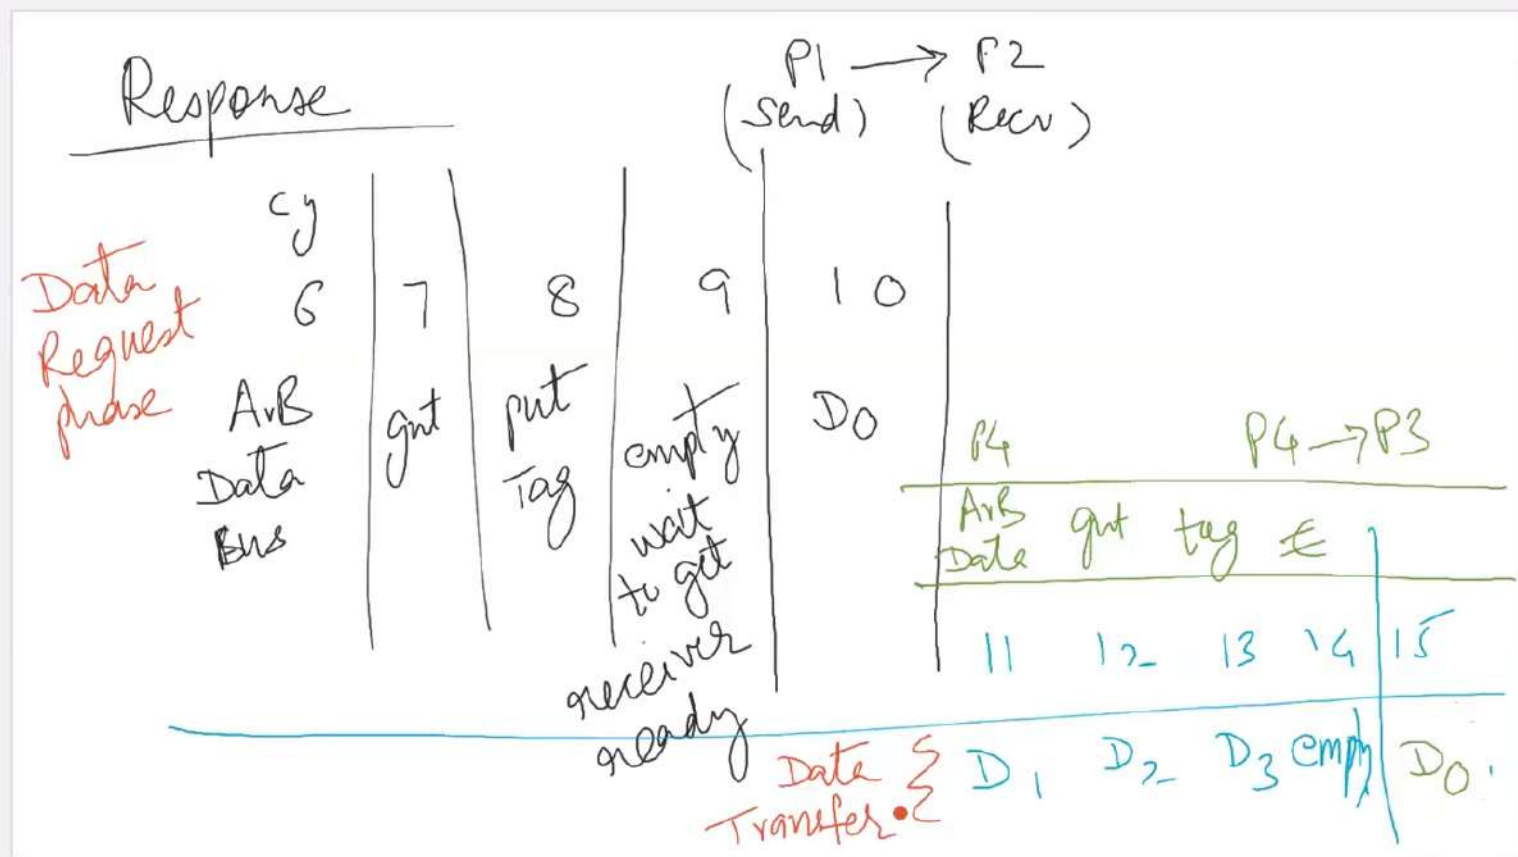
If target is not ready to receive data, it will raise a signal and data transfer will be stalled

Cache local-state updated when data is received

Snoop results => Shared, Dirty, Snoop-valid (inhibit signal)

If snoop-valid signal is disabled then all wait, once enabled, the correct component: memory or cache can start data transfer





Write back & Bus Upgr

- Write back – send block to memory (BusWB)
 - Consists of only request phase
 - Needs both address and data bus.
 - Arbitrates for both simultaneously
- BusUpgr - only request is sent to get exclusive ownership of block
 - here the same bus controller sends a response to its own cache after getting the bus
 - Indicates that write is **committed** and has been **serialised**

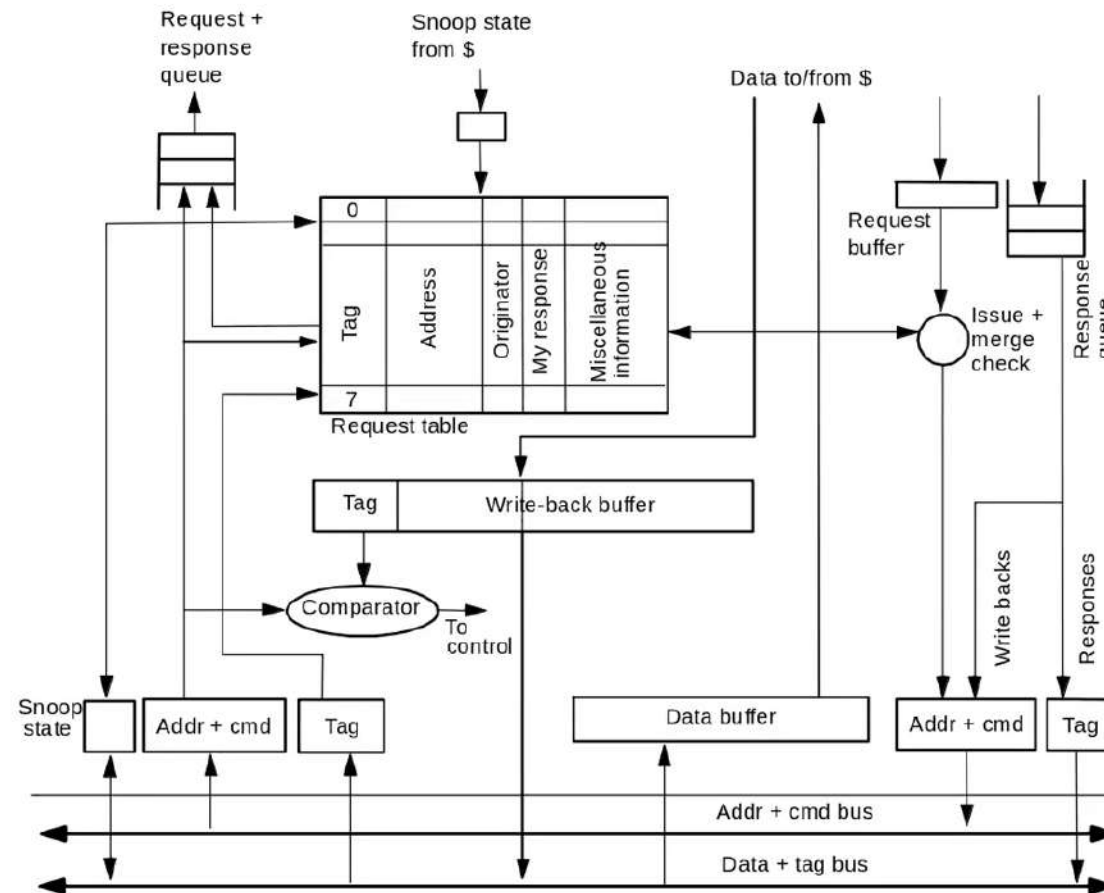


Request table for outstanding requests

- Each cache controller has 8 entry request table
- New requests on bus --> add entry at same index in all request tables
- Entry contains
 - <tag, addr of associated blk, the request type, state of block in local cache, other bits>
 - It is fully associative
 - All entries examined for match
 - Use addr field for request phase and
 - Use tag field in response phase
 - Entry is freed when response is seen on the bus
 - The tag of the freed entry can then be used by newer bus transactions



Bus interface with Request table



Hemangee K. Kapoor

22

+21

AS



SK

SHIVANGI KUMAR

SS

Syam Senkar

SP

SARASWATULA PHANI SAI PRANAV

VA

VARHADE AMEY ANANT



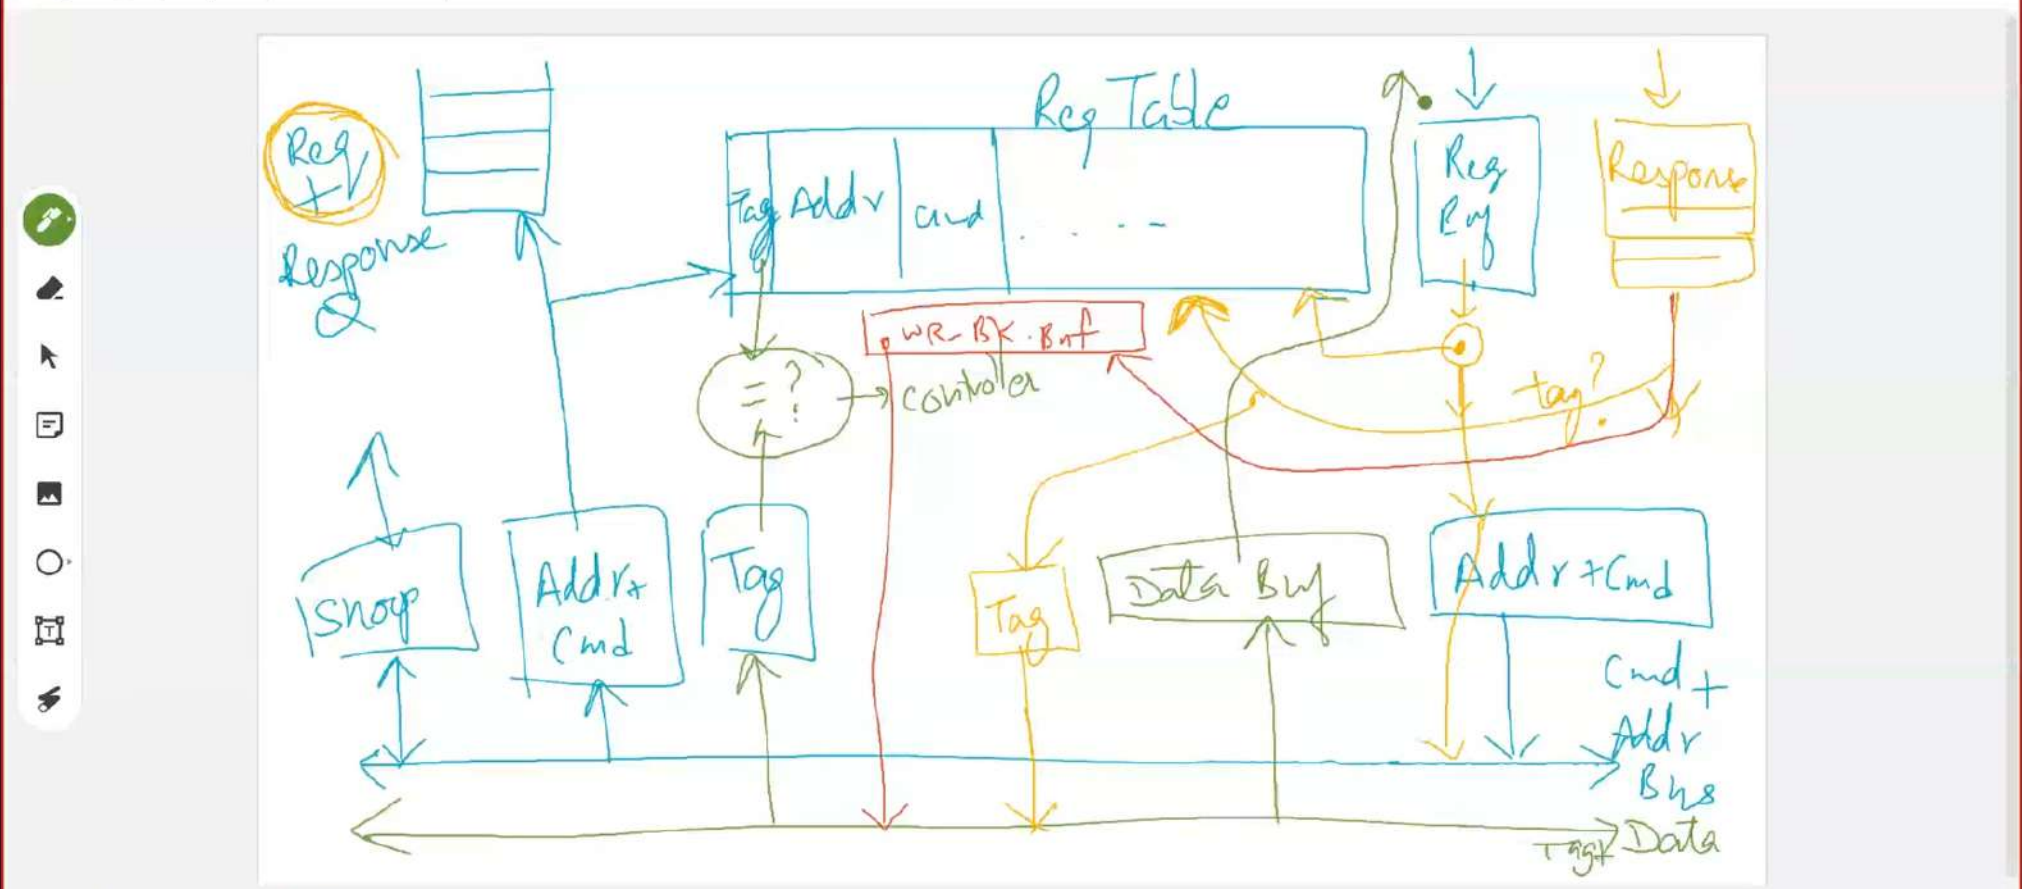
Hemangee Kalpesh Kapoor



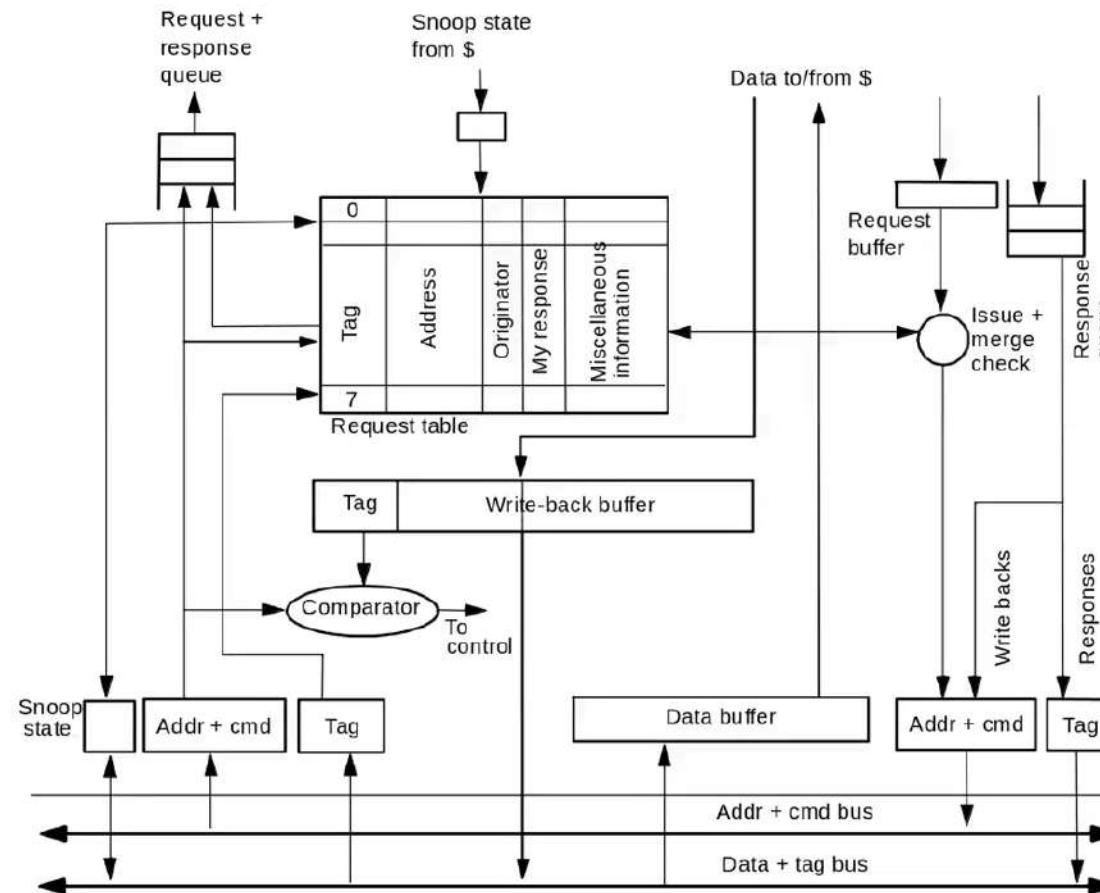
KARTIKEYA SAXENA



Hemangee Kalpesh Kapoor



Bus interface with Request table



Hemangee K. Kapoor

22

+21

AS



SK

SHIVANGI KUMAR

SS

Syam Senkar

SP

SARASWATULA PHANI SAI PRANAV

VA

VARHADE AMEY ANANT



Hemangee Kalpesh Kapoor



KARTIKEYA SAXENA



Hemangee Kalpesh Kapoor

Snoop results and conflicting requests

- Uses variable delay snooping
- 3-wired OR signals: Shared, Dirty, Inhibit
- Snoop results presented when response appears
 - At end of address phase, **determine the response and keep in request table entry**
 - It is also determined **who will respond**. But may take many cycles for data to be ready
 - Write-backs and upgrades don't have data response or snoop result
- Therefore, to match snoop results with their request, the snoop results are presented on the bus by all controllers at the time the actual response is put on the bus
- Controller has record of pending requests in the table
- **Avoid conflicting requests on bus**
 - No new requests are issued for a block that has an outstanding transaction, i.e. entry in the request table
- Remember? : Writes committed when request gets the bus



Flow Control

- Where? At cache controller
 - i.e. At which buffers and who does the control
 - Incoming requests from bus to cache controller are buffered in request queue
 - Cache sub-system sends response in the response buffer, each entry has addr+blk-data. Size is to be limited
 - The cache controller provides flow control by limiting the number of outstanding requests
- Main memory flow control ..



Flow control..

- Main memory flow control
 - 8 outstanding bus transactions are allowed
 - If each generates a write-back, then memory should have enough space
 - Wr-back happen in quick succession as no response phase
- SGI Challenge: both data and addr-bus has NACK line
 - If buffers are full the memory can raise NACK before the request/response transaction reaches the ACK-cycle
 - If NACK=1 then the transaction is cancelled and later retried
 - Back-off and priorities are used to reduce traffic and starvation
- SUN Enterprise: destination (in our case memory) initiates retry when it has free buffer space
 - Source keeps watch for this retry
 - Guaranteed space will still be there, so need only one retry



Path of a Cache Miss

- (1) Handling a Read Miss
- (2) Handling a Write Miss



Handling a Read Miss

- Processor wants to read a block which is not in the cache.
 - Need to issue **BusRd**
- Check if Request table has request for same block
 - Yes: see if you can get the block when the existing request is being served
 - No: issue request and watch-out for race conditions



Handling read miss

Prior request
BusRd

Yes: prior request exists for same block and request is also BusRd

We can also grab data. For this add 2-bits to each entry in the table:

1st bit : "want to grab response"

2nd bit : "I am original requestor"

For this cache the bits are 1st, 2nd = $\langle 1, 0 \rangle$

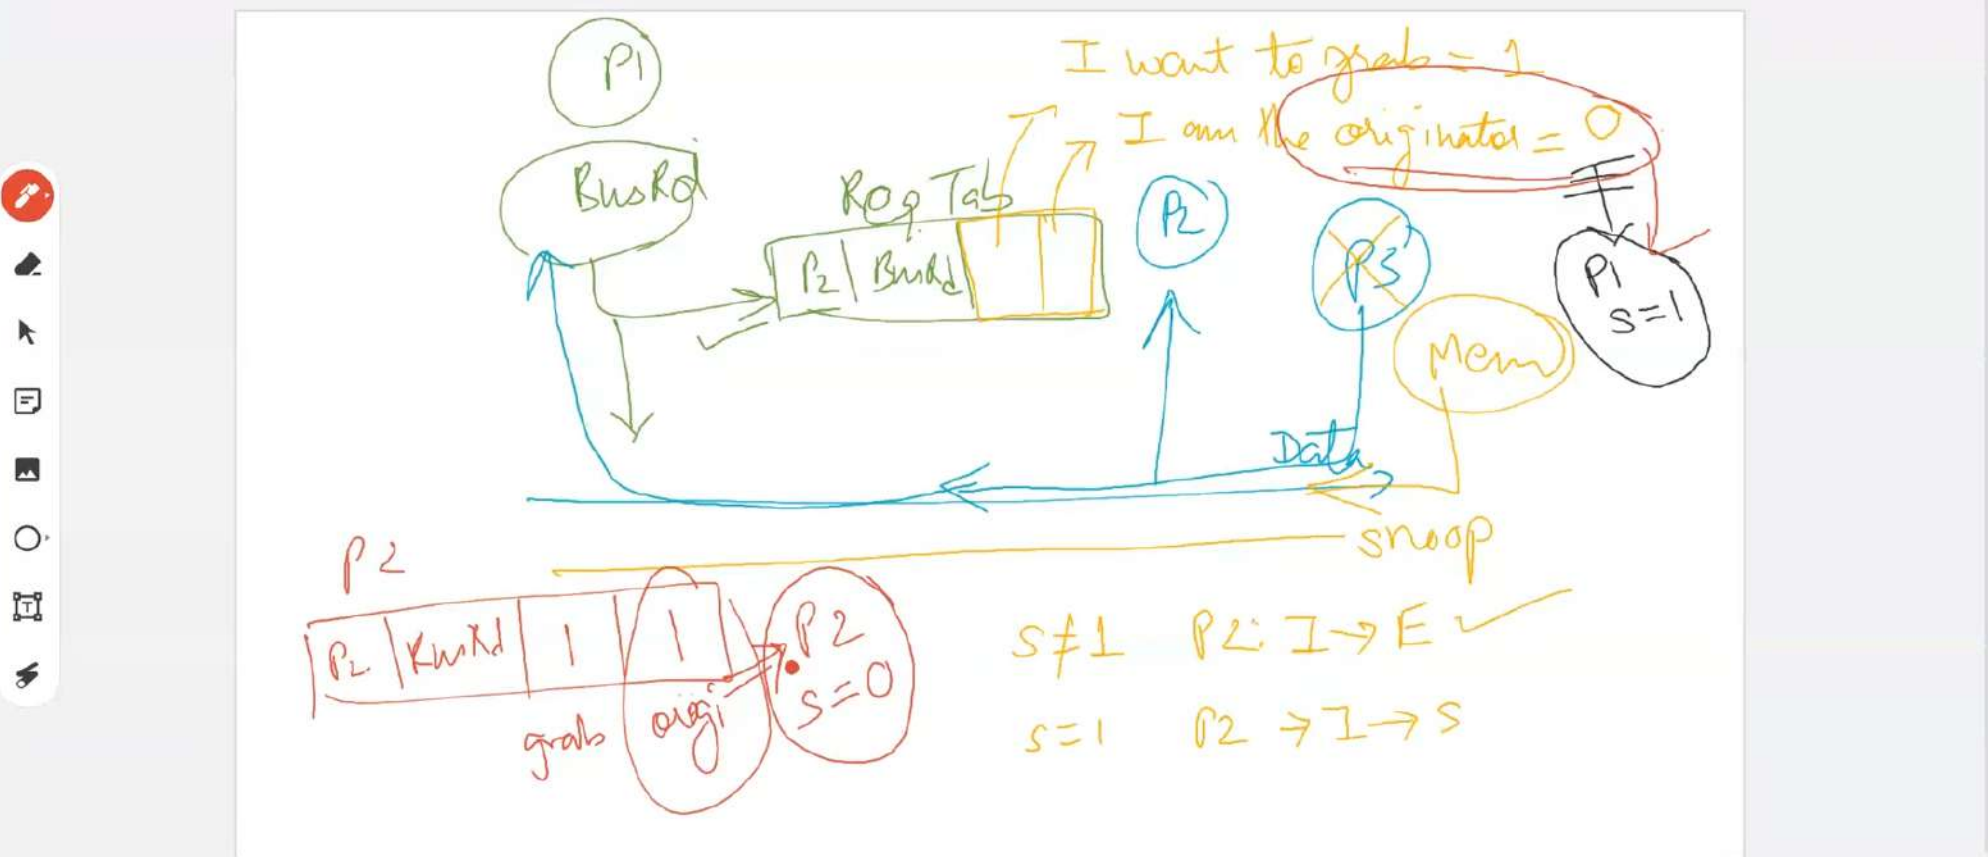
Non-original grabber must assert sharing line so that others will load in 'S' state rather than 'E' state

Original requester will have them as? $\langle 1, 1 \rangle$

Prior request
BusRdX

If incompatible prior request, then wait for the earlier request to complete and later retry





Handling read miss

Prior request
BusRd

Yes: prior request exists for same block and request is also BusRd

We can also grab data. For this add 2-bits to each entry in the table:

1st bit : “want to grab response”

2nd bit : “I am original requestor”

For this cache the bits are 1st, 2nd = $\langle 1, 0 \rangle$

Non-original grabber must assert sharing line so that others will load in ‘S’ state rather than ‘E’ state

Original requester will have them as? $\langle 1, 1 \rangle$

Prior request
BusRdX

-
If incompatible prior request, then wait for the earlier request to complete and later retry



Handling read miss

No: there is no prior request for this block in the request table

Issue request and watch-out for race conditions

Possible race condition=

- (i) controller checks req-table. No conflicting requests
- (ii) before it gets the bus, another conflicting req is granted the bus
- (iv) this cache granted bus immediately next

There is a possibility of 2 conflicting requests on bus

Therefore this cache must check req-table entry just before its own slot on the bus

If it finds an entry then

- (i) issues null-request (no-action) on bus to occupy the given slot and
- (ii) withdraw from further arbitration until the conflicting request is completed



Effect on other processors

- Processor finally is successful in issuing BusRd. What should other cache-controller and memory do?
- Request is entered in the req-table of all cache controllers
- Cache controllers start checking their caches for the block
- Memory does not know if the block is dirty in some cache. So memory also starts fetching the block
- 3 scenarios:
 - (i) Cache has dirty block: Cache wins bus
 - (ii) Cache has dirty block: memory wins bus
 - (iii) Caches do not have dirty block: memory supplies data

