

Directory based Cache Coherence

Topic-4
Chapter-8

Hemangee K. Kapoor

1



ASWATHY N S



SHIVANGI KUMAR



KARTIKEYA SAXENA



SUBRATA ROY



NISHANK SIDDHARTH



RATHOD SAINATH



Hemangee Kaipesh Kapoor

Protocol Optimisations

Hemangee K. Kapoor

2



ASWATHY N S



SHIVANGI KUMAR



KARTIKEYA SAXENA



SUBRATA ROY



NISHANK SIDDHARTH



RATHOD SAINATH



Hemangee Kaipesh Kapoor

Performance

- Performance
 - Network transactions on which cache coherence protocols are built differ from message passing ones
 - (i) they are automatically generated by the communication assist or controllers
 - (ii) individually small in size, only req, response, ack or data
- Each network transaction incurs overhead at requesting processor, communication assist at end-points, network delay
- Processor involved only at requestor. Processor at dirty node or home node is not involved: here CA does the job
- Performance can be improved by
 - (i) protocol optimisation
 - (ii) high-level machine organisation
 - (iii) hardware specialisation to improve communication parameters



ASWATHY N S



SHIVANGI KUMAR



KARTIKEYA SAXENA



SUBRATA ROY



NISHANK SIDDHARTH



RATHOD SAINATH




Hemangee Kaipesh Kapoor

Performance

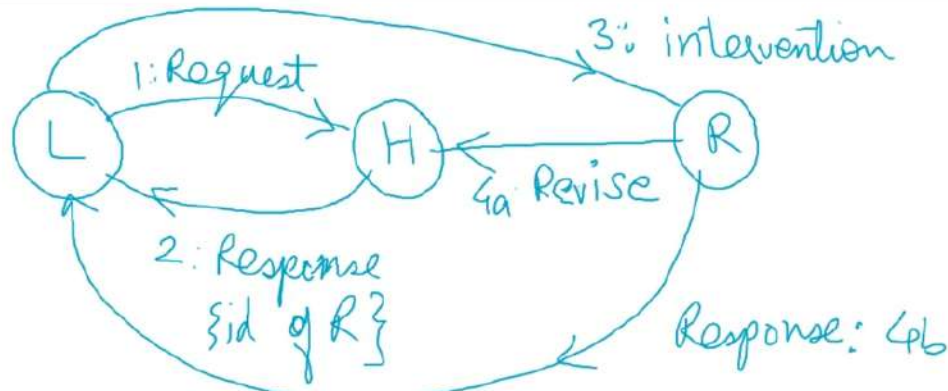
- Performance can be improved by
 - (i) protocol optimisation
 - We will see 3 variants ...
 - (ii) high-level machine organisation
 - mostly 2-level hierarchy with multiprocessor nodes
 - Each node can be snooping or directory
 - Snooping nodes have longer latency but still used
 - (iii) hardware specialisation to improve communication parameters
 - Keep assist occupancy small, tighter integration with memory unit, allow assist to start new transaction while data of older transaction is being fetched, make the assist more specialised
 - Use SRAM directory to reduce look-up time
 - a single bit can be kept with memory to keep track of whether the block is clean or not (so that communication assist is not invoked on a read miss)
 - If assist occupancy is high then pipeline its design or overlap its actions



Protocol optimisations

- Protocol optimisation goals
 - Reduce **#network transactions**
 - Reduces bandwidth demand on the network and the communication assist
 - Reduce **#transactions on critical path** of processor 
 - Can be done by overlapping the transactions needed for a memory operation as much as possible
- Manner in which directory information **stored**, determines the number of network transactions in the critical path of processor
 - Memory based: can be overlapped
 - Cache based: cannot be overlapped
 - In both cases options of improvement are possible
- **Methods of communication**
 - (i) strict Request-Response
 - (ii) intervention forwarding
 - (iii) reply forwarding





STRICT Reg-Response


L = local node

H = home

R = remote


$n/w T_{xh} = 5$

critical path = 4



n/w Txn = 4
cri-path = 4


INTERVENTION
FWD




4
3

REPLY FWD


Protocol enhancements for latency



(a) Strict request-reply



(a) Intervention forwarding



(a) Reply forwarding

L = local requesting node

H = home node

R = remote, dirty node

Strict Req-Response

- 1: request
- 2: carriers owner indentity
- 3: intervention
- 4a: sends data
- 4b: sends data to update “home” and change directory-state
- => we have 4 network transactions in critical path of Read operation and total 5 network transactions
- Reduce this using (2) intervention forwarding ..



Intervention forwarding

- 1: request, 2: intervention, 3: response, 4: response
- Home does not respond to requestor
- But forwards request as intervention to owner asking for data
- Owner sends data to home. Updates state. Home sends data to requestor
- *“Intervention is like a request but is issued in reaction to a request and is directed at a cache rather than memory” [similar to inv but seeks data]*
- Total network transactions reduced from 5 to 4
 - Bandwidth requirement reduced
 - #transactions in critical path is still 4
- So need more aggressive method = Reply forwarding



Intervention forwarding

- 1: request, 2: intervention, 3: response, 4: response
- Home does not respond to requestor
- But forwards request as intervention to owner asking for data
- Owner sends data to home. Updates state. Home sends data to requestor
- *“Intervention is like a request but is issued in reaction to a request and is directed at a cache rather than memory” [similar to inv but seeks data]*
- Total network transactions reduced from 5 to 4
 - Bandwidth requirement reduced
 - #transactions in critical path is still 4
- So need more aggressive method = Reply forwarding



Reply forwarding

- 1: request, 2: intervention: 3a: revision, 3b: response
 - Home forwards the intervention message to owner but here intervention message contains identity of requestor
 - Owner replies directly to requestor
 - Owner also sends revision message to home to update memory data and directory state
- └
- But 3a is not in critical path of Read miss
 - Total #transactions is 4 and #transactions in critical path is 3
 - Called three-message miss




Which to use?



- Type (2) and (3): intervention forwarding and reply forwarding are **NOT strict request-response**
- Since a request at home generates another request (to owner) which in turn generates response. This complicates **deadlock** avoidance
- Intervention forwarding is intermediate in its **latency** and **traffic** characteristics
- At the same time has disadvantage that **all outstanding interventions** are kept track-of at home node
- **Home** has to keep track of **k*P requests** as there can be 'k' requests from each node in a 'P' node system
- **Requestor** has to keep track of **only 'k'** outstanding requests
- Reply forwarding does not require home to keep track and has better performance
- Therefore systems prefer to use reply forwarding



Reducing latency in flat, cache-based protocol

- Scenario: write request. Inv to be sent to i sharers
- Strict request-response
 - In each 'ack' identity of next sharer is sent to home
 - Total #transactions is $2s$: $s = \text{\#sharers}$
 - All $2s$ are in critical path






n/w Txn = 2s
 # cri path = s+1

n/w -
 → s+1
 → s+1.

Intervention fwd and reply fwd


- Intervention fwd

- Total transactions = $2s$
- In critical path = $s+1$
- As each s_i sends inv to next sharer and in parallel sends ack to home



- Reply fwd

- Only last sharer sends ack to home
- Total transactions = $s+1$
- In critical path = $s+1$
- These two types are NOT strict-request-response cases



Correctness

- Similar to snoop, here correctness concerns are
- (1) Protocol ensures relevant blocks are invalidated and data retrieved, state transitions happen (we assume this holds, we will not prove it here)
- (2) serialisation and ordering relationships defined by coherence and consistency model must be preserved
- (3) protocol implementation must be free from deadlock, livelock and (ideally) starvation
- Points (2) and (3) complicated by scalable protocols because
 - (i) multiple copies of block but no single agent that sees all relevant transactions and serialises them
 - (ii) with many processors, large number of requests may be directed towards a single node, creating input buffer problems

Correctness

- Scalable systems have **high latency** which aggravate the problem and makes us use protocol optimisations (As shown earlier)
- Optimisations allow
 - (i) more transactions to be in **parallel** simultaneously, and
 - (ii) do **not preserve** a strict request-response nature, thus complicating correctness

0

Serialisation to a location for coherence


- A given processor must be able to construct a serial order out of all the operations to a given location
 - At least out of all write operations and its own read operation AND
- All processors must see the writes to a given location as having happened in the same order
- For serialisation we need one entity that sees the necessary operations to a given location from all processors and determines the serial order
- (i) easy in BUS based, as bus is the entity
- (ii) in distributed system that does not cache shared data: memory at home can be entity
 - Order in which writes reach home is the order visible to all processors
 - Which write's value a read sees is determined by when the read reaches home memory
- (iii) in distributed system with coherent caching ..

Serialisation to a location for coherence

- (iii) in distributed system with coherent caching
 - **Home** memory is likely candidate as serialising entity, atleast in flat-directory
 - (a) if home can satisfy all requests itself then it processes them in **FIFO** order and determines serialisation
 - But with multiple copies, visibility at home does not imply visibility by all processors
 - Easy to generate scenario when different processors see different orders than what are seen by home
 - Examples: (1) update protocol (2) inv protocol write request

Ex: see diff orders: update prot

- (1)
Update based protocol and network does not preserve **point-to-point order** of transactions between the end-points
- If two write requests for shared data arrive at home **in one order**, the updates they generate may arrive at the copies in **different orders**




H: upd1 ; upd2

N3: upd1 ; upd2

N4: upd2 ; upd1

Ex: see diff orders: inv protocol

- (2) in inv-based protocol, block in dirty state in a node.
Two other nodes issue a **Read exclusive** request
 - In strict request-response protocol, the home will give identity of owner to both requestors
 - Note that the Home node will not get updated with new owner until the revision message comes from owner. In the mean time a new requestor can come and Home will give the same owner ID to this new requestor.
 - Then the requestor will go to owner to get data
 - Order in which requestor reached home and order in which they reach owner may be different
 - Which entity provides globally consistent serialisation in this case?



H $L1 \rightarrow L2$
R $L1 \rightarrow L2$
 $L2 \rightarrow L1$

Ex: see diff orders: inv protocol

- (2) in inv-based protocol, block in dirty state in a node.
Two other nodes issue a Read exclusive request
 - In strict request-response protocol, the home will give identity of owner to both requestors
 - Note that the Home node will not get updated with new owner until the revision message comes from owner. In the mean time a new requestor can come and Home will give the same owner ID to this new requestor.
 - Then the requestor will go to owner to get data
 - Order in which requestor reached home and order in which they reach owner may be different
 - Which entity provides globally consistent serialisation in this case?