

Memory Consistency

Hemangee K. Kapoor

1

NS

NALABOLU SAN...

YK

YOGESH KUMAR



DODDAVULA LIK...

AS

ASWATHY N S

KS

KUSHAL SANGW...

SU

SWATI UPADHYAY

DS

DRISTIRON SAIKIA



IMJUNGLA LON...



K CHITRA



TANVISH

DN

DARSHIT NAGAR

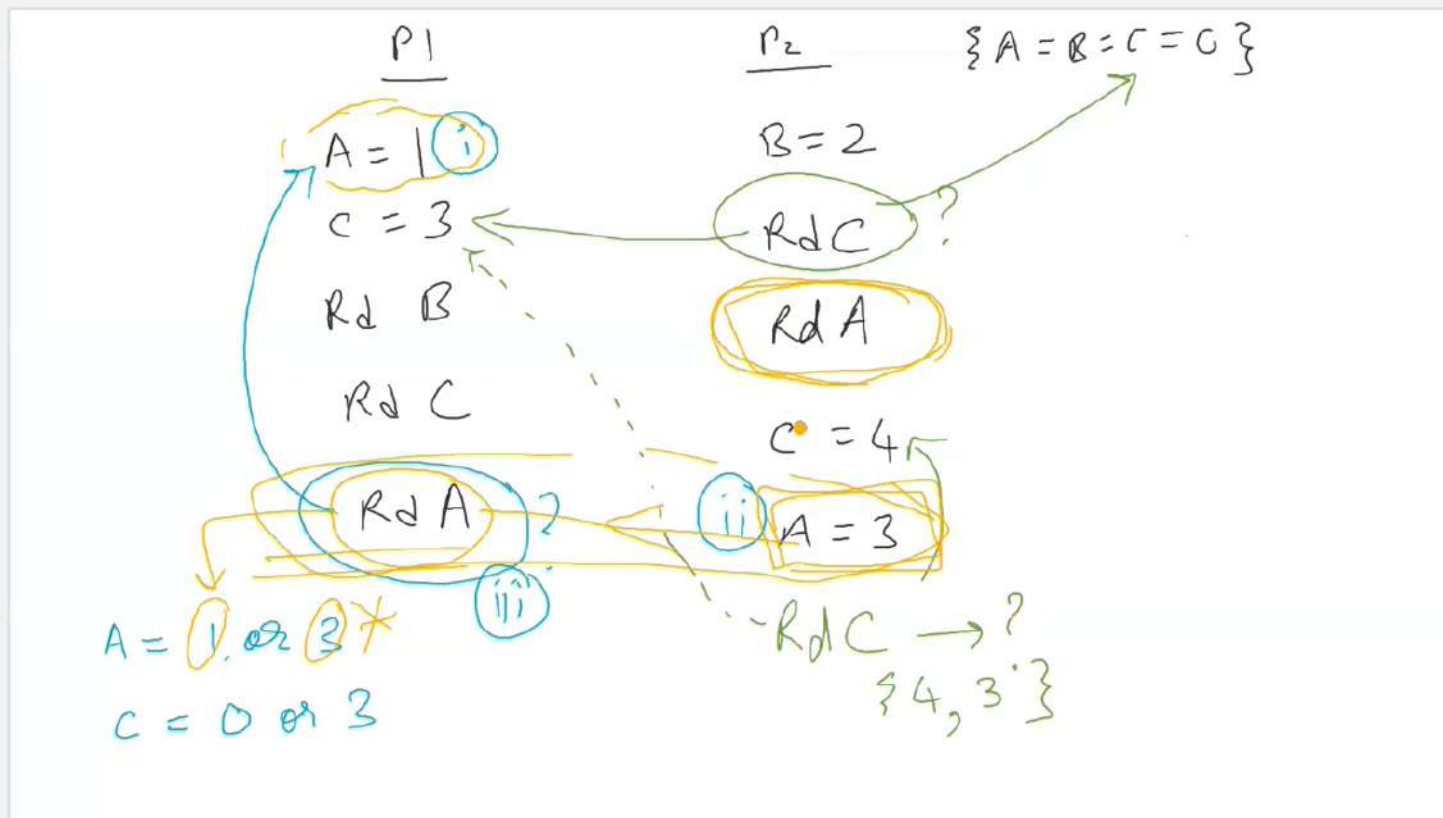


Hemangee Kalpe...



Hemangee Kalpe...





TANVISH

NS

YK

VOGESH KUMAR

DODDAVULA LIK...

AS

ASWATHY N S

KS

KUSHAL SANGW...

SU

SWATI UPADHYAY

DS

DRISTIRON SAKIA

IMLIJUNGLA LON...

K CHITRA

DN

DARSHIT NAGAR

Hemangee Kalpe...

AM

KOUSIK RAJESH

ANSHUL MITTAL

A

AMIT

SP

SARASWATULA P...

SUDRATA ROY

+2

28-Oct-2021 - Google x 28_Oct - Mentimeter +

jamboard.google.com/d/1eX86cCvyxTvQIE1-JeVOiBTLoI4eHZ8Em4fipwDErs/viewer?f=2

28-Oct-2021

3/3

Set background Clear frame

$\{A=0\}$

P1

$A=1$

$flag=1$

P2

$while(flag == 0);$

$\rightarrow print A \rightarrow 0$

$\times 1$ or

TANVISH

NS

NALABOLU SAN...

YK

VOGESH KUMAR

DODDAVULA LIK...

AS

ASWATHY N S

KS

KUSHAL SANGW...

SU

SWATI UPADHYAY

DS

DRISTIRON SAKIA

IMLIJUNGLA LON...

K CHITRA

DN

DARSHIT NAGAR

Hemangee Kalpe...

AM

KOUSIK RAJESH

ANSHUL MITTAL

A

AMIT

SP

SARASWATULA P...

SUDRATA ROY

+2

Coherence

- Writes to a location become visible to all in the same order
- But when does a write become visible?
- Often we want to ensure that a read returns the value of a particular write, i.e. we want to establish an order between a write and a read
- Typically use event synchronisation, i.e. use more than one location

• Ex:

P1	P2
<i>// Assume initially A=0 ; and flag=0</i>	
A = 1;	while(flag == 0); // spin idly
flag = 1;	print A; // i.e. read A

- Programmer intuition => 'print A' result is '1'
- We use access to another location ('flag') to preserve the desired order for same location ('A')



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAWULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUDRATA ROY



+2

Example ...

- We assume that write-A by P1 becomes visible to P2 before write-flag
- AND
- Read of flag by P2 which breaks while-loop completes before read-A in P2
- These are program orders within P1 and P2 but to different locations => which is not implied by coherence
- Since coherence only requires that new value of A (P1) becomes eventually visible to P2 and not necessarily before the new value of flag is observed



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAWULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT




SARASWATULA P...



SUDRATA ROY




+2



Participant list:

- TANVISH
- NS
- YK
- VOGESH KUMAR
- DODDAVULA LIK...
- AS
- ASWATHY N S
- KS
- KUSHAL SANGW...
- SU
- SWATI UPADHYAY
- DS
- DRISTIRON SAKIA
- IMLIJUNGLA LON...
- K CHITRA
- DN
- DARSHIT NAGAR
- Hemangee Kalpe...
- KOUSIK RAJESH
- AM
- ANSHUL MITTAL
- A
- AMIT
- SP
- SARASWATULA P...
- SUDRATA ROY
- +2



Participant list:

- TANVISH
- NS
- YK
- VOGESH KUMAR
- DODDAVULA LIK...
- AS
- ASWATHY N S
- KS
- KUSHAL SANGW...
- SU
- SWATI UPADHYAY
- DS
- DRISTIRON SAKIA
- IMLIJUNGLA LON...
- K CHITRA
- DN
- DARSHIT NAGAR
- Hemangee Kalpe...
- AM
- KOUSIK RAJESH
- ANSHUL MITTAL
- A
- AMIT
- SP
- SARASWATULA P...
- SUDRATA ROY
- +2

Another example of ordering

P1	P2
// initially A = B = 0;	
(1a) A = 1;	(2a) print B;
(1b) B = 2;	(2b) print A;

- Here programmers intuition is not clear, as there are no synchronisation variables
- Should we assume that if print-B = 2 then print-A=1?
- Whatever is the answer, but coherence says nothing about this



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAVULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUDRATA ROY



+2

What is needed?

- We need an ordering model for clear semantics
 - Across different locations as well
 - So that programmers can reason about what results are possible
- This is the memory consistency model for shared address space
- Specifies constraints on the order in which memory operations must appear to be performed (i.e. to become visible to the processors) wrt one another
- This includes operations to the same location, or different location and by the same process or different processes
- Memory consistency subsumes coherence



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAVULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUDRATA ROY



+2



Why Consistency is Important?

- Consistency model defines correct behaviour
- It is a contract between the system and programmer
- Coherence is not visible to software (i.e. not architectural)
 - Therefore it enables new system to present same consistency model despite using newer (fancier) coherence protocols
- Consistency model restricts ordering of load/store and does not care about ordering of coherence messages

Hemangee K. Kapoor

6



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAVULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATI P...



SUDRATA ROY




+2



Sequential Consistency

- Informally: A desirable ordering for shared address space: the reasoning that allows a multi-threaded program to work on a uni-processor (using any interleaving) should hold for the multi-threaded program running on more than 1 processor
- Leslie Lamport 1979

"A multiprocessor is sequentially consistent if the result of any execution is the same as if the operations of all the processors were executed in some sequential order, and the operations of each individual processor appear in this sequence in the order specified by its program"
- Abstraction = a multi-tasking uniprocessor



Hemangee K. Kapoor

7



TANVISH

NS
NALABOLU SAN...YK
YOGESH KUMAR

DODDAWULA LIK...

AS
ASWATHY N SKS
KUSHAL SANGW...SU
SWATI UPADHYAYDS
DRISTIRON SAIKIA

IMLIJUNGLA LON...



K CHITRA

DN
DARSHIT NAGAR

Hemangee Kalpe...



KOUSIK RAJESH

AM
ANSHUL MITTALA
AMITSP
SARASWATULA P...

SUDRATA ROY



+2



Sequential consistency

- Ordering of data accesses within a process are in program order and across processes they are some interleaving of the program orders
- In the memory model
 - Each process issues memory operations and completes **one at a time**, **atomically** and **in program order**
 - [One-at-a-time] = A memory operation does not appear to be issued until the previous one from that process has completed
 - [atomically] = operations in the interleaved order are atomic, i.e. it should appear **globally** as if one operation in this order executes and completes before the next one begins
 - **It is not important in what order the memory operations actually issue or even complete. What matters for sequential consistency is that they appear to complete in a manner that satisfies the constraints described above**
 - Program order = as seen by the programmer (and not as optimised by compiler, etc.)



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAVULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUDRATA ROY



+2

Sequential consistency

- Ordering of data accesses within a process are in program order and across processes they are some interleaving of the program orders
- In the memory model
 - Each process issues memory operations and completes **one at a time**, **atomically** and **in program order**
 - [One-at-a-time] = A memory operation does not appear to be issued until the previous one from that process has completed
 - [atomically] = operations in the interleaved order are atomic, i.e. it should appear **globally** as if one operation in this order executes and completes before the next one begins
 - **It is not important in what order the memory operations actually issue or even complete. What matters for sequential consistency is that they appear to complete in a manner that satisfies the constraints described above**
 - Program order = as seen by the programmer (and not as optimised by compiler, etc.)



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAVULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUDRATA ROY



+2



SC: example

- What matters is order in which operations appear to execute, not the chronological order of events

P1	P2
// initially A = B = 0;	
(1a) A = 1;	(2a) print B;
(1b) B = 2;	(2b) print A;

- Possible outcomes: (A,B) : (0,0), (1,0), (1,2)
- What about (0,2)?
- Program order = 1a->1b and 2a → 2b

- A=0 => 2b -> 1a, which implies 2a -> 1b
- B=2 => 1b -> 2a, which leads to contradiction

- What is actual execution 1b->1a->2b->2a
- This execution sequence appears just like: 1a-> 1b-> 2a -> 2b : as visible from results (A,B) = (1,2)

- However the sequence: 1b-> 2a-> 2b->1a is not SC since it would produce result (0,2)
- i.e. it does not matter that they actually execute and complete out-of-order; as long as results of execution show the same effect to have executed in-program-order



TANVISH



NALABOLU SAN...



YOGESH KUMAR



DODDAWULA LIK...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



DRISTIRON SAIKIA



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUDRATA ROY



28-Oct-2021 - Google | 28_Oct - Mentimeter

jamboard.google.com/d/1eX86cCvyxTvQIE1-JeVOiBTLoI4eHZ8Em4fipwDErs/viewer?f=3

28-Oct-2021

4/4

Set background Clear frame

Share

P1
 1a) A = 1
 1b) B = 2

P2
 print B 2a)
 print A 2b)

o/p
 [A, B]
 0, 0 ✓
 1, 0 ✓
 1, 2 ✓

is o/p $\langle 0, 2 \rangle$ possible?

po

1a → 1b	2b → 1a → 1b
2a → 2b	1b → 2a → 2b

2b → 1a → 1b → 2a → 2b

TANVISH
 NS
 NALABOLU SAN...

YK
 YOGESH KUMAR
 DODDAVULA LIK...

AS
 ASWATHY N S
 KUSHAL SANGW...

SU
 SWATI UPADHYAY
 DRISTIRON SAKIA

IMLIJUNGLA LON...
 K CHITRA

DN
 DARSHIT NAGAR
 Hemangee Kalpe...

AM
 KOUSIK RAJESH
 ANSHUL MITTAL

A
 AMIT
 SP
 SARASWATULA P...

+2
 SUDRATA ROY

28-Oct-2021 - Google x 28_Oct - Mentimeter +

jamboard.google.com/d/1eX86cCvyxTvQIE1-JeVOiBTLoI4eHZe8Em4fipwDErs/viewer?f=4

28-Oct-2021

Set background Clear frame

5/5

Share

Handwritten notes on a Jamboard:

$$1b \rightarrow 1a \rightarrow 2a \rightarrow 2b \quad o/p = \{1, 2\} \checkmark \underline{\underline{sc}}$$

POX (Process Order X)

$$1a \rightarrow 2b \rightarrow 2a \rightarrow 2b$$

$$1a \rightarrow 2a \rightarrow 2b \rightarrow 1b \rightarrow \{1, 0\}$$

$$A=1 \quad B=0 \quad A=1 \quad B=2$$


$$1b \rightarrow 2a \rightarrow 2b \rightarrow 1a \rightarrow \{0, 2\} \underline{\underline{not sc}}$$

Participants:

- TANVISH
- NS
- YK
- VOGESH KUMAR
- DODDAVULA LIK...
- AS
- ASWATHY N S
- KS
- KUSHAL SANGW...
- SU
- SWATI UPADHYAY
- DS
- DRISTIRON SAKIA
- IMLIJUNGLA LON...
- K CHITRA
- DN
- DARSHIT NAGAR
- Hemangee Kalpe...
- AM
- KOUSIK RAJESH
- ANSHUL MITTAL
- A
- AMIT
- SP
- SARASWATULA P...
- SUDRATA ROY
- +2

Implementing SC

- Implementing SC requires the system to satisfy two constraints (1) program order, (2) atomicity
- (1) Program order
 - Memory operations issued by a process must appear to execute (become visible to others + itself) in program order
- (2) Atomicity
 - In the overall hypothetical total order, one memory operation should appear to complete wrt all processors before the next-one in the total order is issued
 - Guarantees that total order is consistent across processes
- Tricky point is making writes atomic
 - When there are multiple copies across the system that need invalidation/updating
 - It ensures that:



Then effect of wr-B is not seen by P1
or other process until all others have seen wr-A



TANVISH



YOGESH KUMAR



NALABOLU SAN...



ASWATHYN S



KUSHAL SANGW...



SWATI UPADHYAY



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUBRATA ROY



RATHOD SAINATH



NAREESH BHARAS...



Hemangee Kalpe...

Write atomicity

- Write atomicity extends write-serialisation: in that wr-serialisation (coherence) is wrt same location, while wr-atomicity caters to all locations
- All see writes in same order
 - To same location = wr-serialisation
 - To all locations = wr-atomicity

Ex:

P1	P2	P3
<code>A = 1;</code>	<code>while (A == 0);</code>	
	<code>B = 1;</code>	<code>while (B == 0);</code>
		<code>print A;</code>

- P2 waits for A to become 1 then makes B=1
- P3 waits for B to become 1 then prints A
- Due to transitivity print-A must be 1
- But if P2 is allowed to go on past the read-A and do write-B before P3 has seen the new value of A, then P3 reads new-B and old-A (from its cache) violating sequentially consistent intuition
 - Here P2 should not be allowed to go past Read-A until P3 has also seen the new value of A



TANVISH



YOGESH KUMAR



NALABOLU SAN...



ASWATHYN S



KUSHAL SANGW...



SWATI UPADHYAY



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



KOUSIK RAJESH



ANSHUL MITTAL



AMIT



SARASWATULA P...



SUBRATA ROY



RATHOD SAINATH



NAREESH BHARAS...



Hemangee Kalpe...



Write atomicity

- Write atomicity extends write-serialisation: in that wr-serialisation (coherence) is wrt same location, while wr-atomicity caters to all locations
- All see writes in same order
 - To same location = wr-serialisation
 - To all locations = wr-atomicity

Ex:

P1	P2	P3
<code>A = 1;</code>	<code>while (A == 0);</code>	
	<code>B = 1;</code>	<code>while (B == 0);</code>
		<code>print A;</code>

- P2 waits for A to become 1 then makes B=1
- P3 waits for B to become 1 then prints A
- Due to transitivity print-A must be 1
- But if P2 is allowed to go on past the read-A and do write-B before P3 has seen the new value of A, then P3 reads new-B and old-A (from its cache) violating sequentially consistent intuition
 - Here P2 should not be allowed to go past Read-A until P3 has also seen the new value of A



TANVISH



YOGESH KUMAR



NALABOLU SAN...



ASWATHYN S



KUSHAL SANGW...



SWATI UPADHYAY



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



ANSHUL MITTAL



SARASWATULA P...



SUBRATA ROY



RATHOD SAINATH



PAIDIMARRI MA...



VADIGE PRANEET...



Sufficient conditions for preserving SC

- (1) Every process issues memory operations in program order
- (2) After a write operation is issued (i.e. leaves processor and is presented to memory system including cache), the issuing process waits for the write to complete (i.e. wrt all processors) before issuing its next operation
- (3) After a read is issued, the issuing process waits for the read to complete, and for the write whose value is being returned by the read to complete, before issuing its next operation
 - i.e. if the write (whose value is being returned) has performed wrt this processor, then the processor must wait until it is performed wrt all processors
- The 3rd condition ensures write atomicity and is quite demanding
 - It is not a simple local constraint, because the read must wait until the logically preceding write has become globally visible
- NOTE: These are sufficient conditions, but more than necessary,
 - i.e. SC can be preserved with less serialisation in many cases

Hemangee K. Kapoor

12



TANVISH



YOGESH KUMAR



NALABOLU SAN...



ASWATHY N S



KUSHAL SANGW...



SWATI UPADHYAY



IMLIJUNGLA LON...



K CHITRA



DARSHIT NAGAR



Hemangee Kalpe...



ANSHUL MITTAL



SARASWATULA P...



SUBRATA ROY



RATHOD SAINATH



PAIDIMARRI MA...



VADIGE PRANEET...

Exit fullscreen (F)

Hemangee Kalpe...

