

eting in "General"

06:06:28 UTC

e Kalpesh

Organised by

Hemangee Kalpesh
Kapoor

Channel

General



Global (Barrier) Event Synchronisation

- Centralised Software Barrier
- Shared counter
 - To maintain number of processes arrived at Barrier
 - Incremented by each arriving process
 - Increment must be mutually exclusive
- Process increments count
 - If $\text{count} == p$ then it is last process
 - Else busy-wait till Barrier flag is ON
- The last process sets Barrier flag and releases all $(p-1)$ waiting processes

Simple Centralised Barrier algo

```

struct bar_type {
    int counter;
    struct lock_type lock;
    int flag = 0;
} bar_name;
BARINIT (bar_name) {
    LOCKINIT(bar_name.lock);
    bar_name.counter = 0;
}
BARRIER (bar_name, P) {
    int my_count;
    LOCK (bar_name.lock);

    if (bar_name.counter == 0) {
        bar_name.flag = 0; /* first one */ }

    my_count = bar_name.counter++;

    UNLOCK (bar_name.lock);

    if (my_count == P) { // last one to arrive
        bar_name.counter = 0; //reset count,
        bar_name.flag = 1; } // set flag
    else {
        while (bar_name.flag == 0); } } // busy wait

```


Centralised Barrier with Sense Reversal

- Problem in above barrier, if barrier operation is use consecutively using same bar variable. EX:

some computation

BARRIER(bar1, p);

Some more computation

BARRIER(bar1, p);

- Sample execution leading to the problem. Steps =

1. P1 || P2 || P3 || P4

2. BARRIER (1st)

3. (P1 last, so releases all. P1 makes counter=0; flag=1;) || (P2, P3, P4 wait for flag to be 1)

4. P2, P3 see flag=1 [[P4 has not seen flag=1 as it may be in waiting queue, i.e. not scheduled by OS]]

5. therefore P1, P2, P3 do more computation and again wait for BARRIER

6. P1, P2, P3 has made flag=0 (Reset) but P4 waits for flag=1 (its old instance)

7. when P1, P2, P3 are done they wait for flag to be '1' which will be done by the last processes (here P4)

=> BUT P4 is waiting for flag to be 1 (from earlier instance of P1)

Therefore BARRIER will never be resolved

Solution

- (1) Prevent processes to enter new instance of barrier until all have exited the previous of same barrier
- (2) Use another counter and do not reset flag in new instance until counter has turned to p
 - This new counter counts processes that leave barrier
 - But having more counters incurs latency and contention
 - Current setup requires to reset flag when count=p
- (3) Better solution: Do not reset flag
 - Make processes to wait for a new value of flag for every instance
 - e.g. wait for flag=1 then wait for flag to become '0' then again '1', etc.
 - Maximum we can have processes in two barriers:
 - One old pending and one new on-going
 - We need flag in 2 senses only: '0' and '1'
 - Therefore called Toggle ==> called sense reversal