

Convergence

- Various ways to support shared memory on MPI machines and MPI on shared memory machines
- Evolution and role of software has blurred the boundary
 - send/recv supported on shared address space machines using buffers
 - Global address space available
- Now the underlying machine structure has converged towards a common organisation + communication assist / infrastructure
- Hardware organisation also converging
 - Tighter integration of network interface with the cores gives low latency and high bandwidth
- Even clusters of workstations/SMPs are parallel systems
 - These depend on fast system area network



Other types of Parallel Architectures

Hemangee K. Kapoor

32

+34

AG



TU



SUBRATA ROY



NISHANK SIDDHARTH

NS

NISHITHA SHARMA



DEVANSHI GUPTA

DG



VARHAIDE AMEY ANANT

VA

AMIT

A

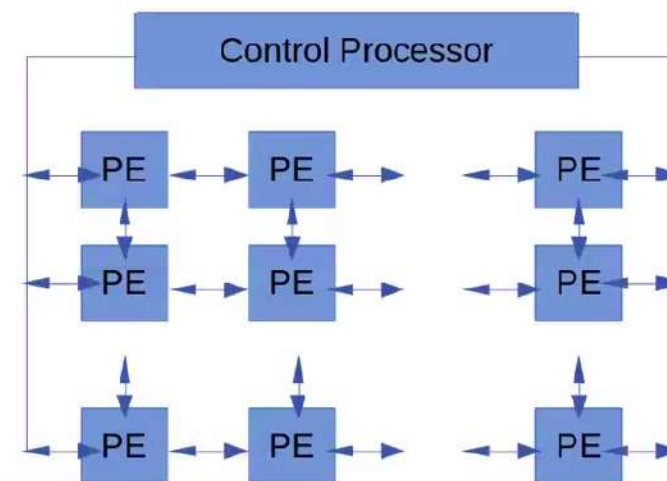


Hemangee Kaipesh Kapoor

Data Parallel Processing

- Here, operations can be performed in parallel on each element of a large regular data structure such as an array or matrix
- Gave rise to Flynn's Taxonomy that categorizes designs in terms of number of distinct instructions issued at a time and the number of data elements they operate on
- SISD: old style processor
- SIMD: like GPUs
- MIMD: shared memory, MPI, multi-programmed etc.
- MISD: no meaning

SIMD = data parallelism
Same instruction on
different data



Hemangee K. Kapoor


33

Vector Processors

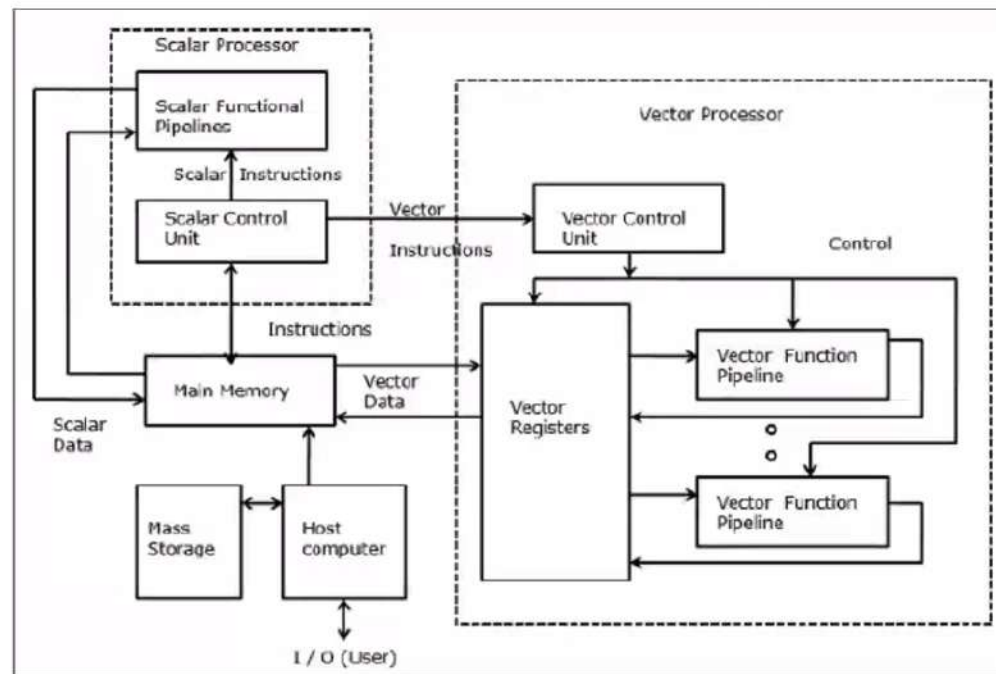
- These operate on arrays or vectors of data
- CPUs operate on individual data elements or scalars
- Recent systems have
 - Vector registers: length of reg=4 to 128, 64-bit elements
 - Vectorised and pipelined functional units
 - Vector instruction
 - Interleaved memory
 - Strided memory access and hardware scatter/gather
 - Access 1st, 5th, 9th element
 - (Scatter = write, Gather = read) at regular intervals (1st, 2nd, 4th, 8th element)



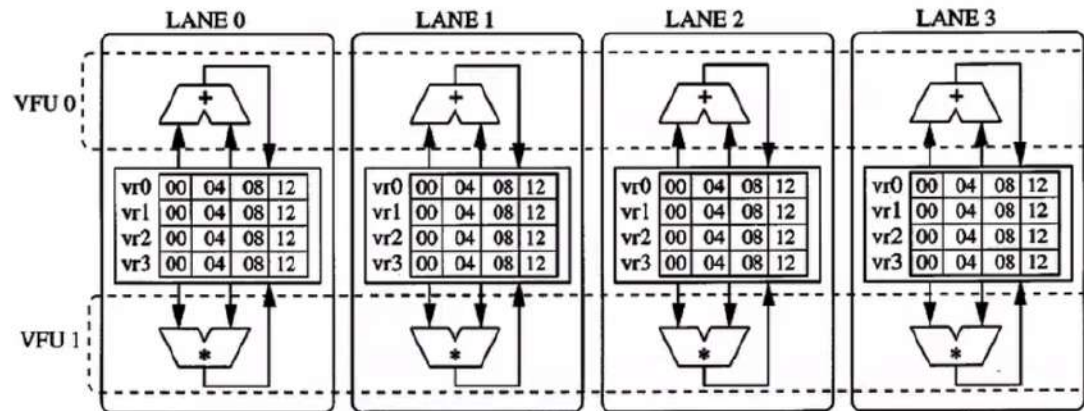
Set background Clear frame



Vector Processors



Vector Processors



C code

```
for (i=0; i<64; i++)
    C[i] = A[i] + B[i];
```

Scalar Code

```
LI R4, 64
loop:
    L.D F0, 0(R1)
    L.D F2, 0(R2)
    ADD.D F4, F2, F0
    S.D F4, 0(R3)
    DADDIU R1, 8
    DADDIU R2, 8
    DADDIU R3, 8
    DSUBIU R4, 1
    BNEZ R4, loop
```

Vector Code

```
LI VLR, 64
LV V1, R1
LV V2, R2
ADDV.D V3, V1, V2
SV V3, R3
```



Vector Processors

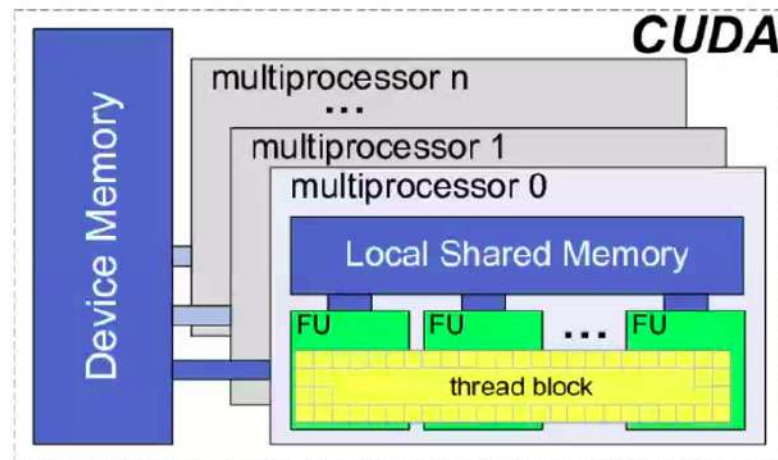

- For matrix multiplication: $A(i,j) = A(i,j) + B(i,k) * C(k,j)$
 - B or C accesses will not be to adjacent or consecutive locations (800 bytes apart)
 - Stride: distance separating the elements that are to be merged into a single vector
 - Caches do unit stride
 - LVWS (load vector with stride) instruction
 - SVWS (store vector with stride)
 - Strides can cause bank conflicts and a stall may occur
- Vector systems have special hardware to accelerate strided access
- Fast and easy to use
- Have high memory bandwidth
- BUT cannot handle irregular data structures and have finite limit on scalability. To handle ever larger problems => larger vectors !



Graphics Processing Units (GPUs)

- Real-time graphics APIs use points, lines and triangles to internally represent the surface of an object
- They use graphics processing pipelines to convert the internal representation into an array of pixels to be sent to the screen
- The working of the graphics pipeline stages is specified by shader functions
- Shader functions are very short (few lines) and several of them can be executed in parallel
- GPUs have large number of ALUs on each GPU core
- Processing single image requires 100s of MB of data

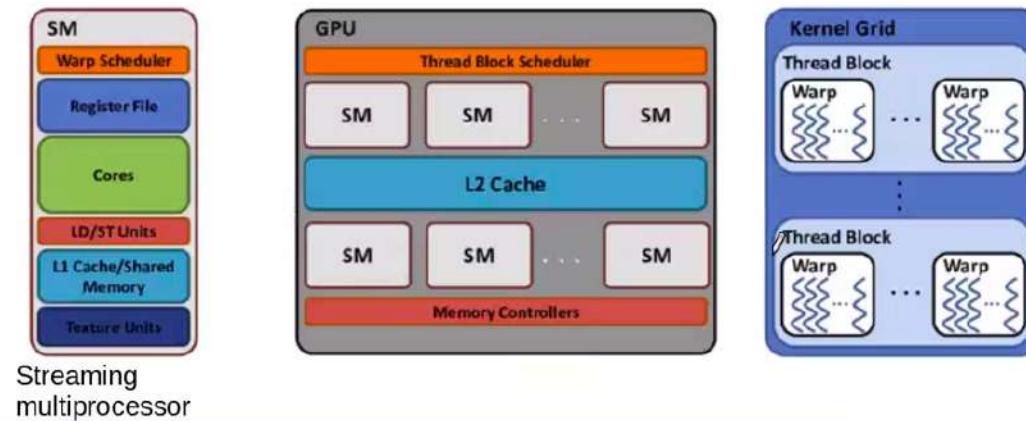




GPUs

- GPUs have
 - High data rates
 - Hardware multi-threading
 - Store state of hundred suspended threads
- We need lots of work to keep the GPUs busy !
- SIMD but some GPUs do multi-program

Other architectures
Are
Dataflow
And
systolic



Streaming
multiprocessor

Hemangee K. Kapoor

40

Shared Memory Multiprocessors

Hemangee K. Kapoor

1

+33

AG



DS



SUBRATA ROY



NISHANK SIDDHARTH

NS

NISHITHA SHARMA

DG

DEVANSHI GUPTA

VA

VARHADE AMEY ANANT

A

AMIT



Hemangee Kaipesh Kapoor

Shared Memory

Hemangee K. Kapoor

2

+33

AG



DS



SUBRATA ROY



NISHANK SIDDHARTH

NS

NISHITHA SHARMA

DG

DEVANSHI GUPTA

VA

VARHADE AMEY ANANT

A

AMIT



Hemangee Kaipesh Kapoor

Prevalent Parallel Architecture

- The most prevalent parallel architecture is multiprocessor of small to moderate scale that provides
 - a global physical address space + symmetric access to all of the main memory from any processor (SMP = Symmetric Multi Processor)
- Each processor has its own cache
- All processors are connected over a common interconnect
- These also form building blocks for larger scale system
- Efficient sharing of resources like CPU + Main Memory enable their use as
 - Throughput Engines
 - Attractive for Parallel Programming
 - Useful for Operating Systems



Prevalent Parallel Architecture

- The most prevalent parallel architecture is multiprocessor of small to moderate scale that provides
 - a global physical address space + symmetric access to all of the main memory from any processor (SMP = Symmetric Multi Processor)
- Each processor has its own cache
- All processors are connected over a common interconnect
- These also form building blocks for larger scale system
- Efficient sharing of resources like CPU + Main Memory enable their use as
 - Throughput Engines
 - Attractive for Parallel Programming
 - Useful for Operating Systems

Hemangee K. Kapoor

3

+33

AG



DS



SUBRATA ROY



NISHANK SIDDHARTH

NS

NISHITHA SHARMA

DG

DEVANSHI GUPTA

VA

VARHADE AMEY ANANT

A

AMIT



Hemangee Kaipesh Kapoor

Why is shared memory useful?

- Throughput engine
 - For multiple sequential jobs with varying memory and CPU requirements
- Useful for parallel programming
 - As they provide ability to access any memory location from any processor using ordinary LOAD/STORE
 - Automatic movement and replication of shared data in the local caches
- Useful for OS
 - As the different processes of OS shared data structures and easily run on different processors



Why is shared memory useful?

- User processes read/write shared virtual address
- Realised as LOAD/STORE on shared physical address
- On the contrary, Message passing needs support from intermediate software layer (i.e. run-time library)
- Message passing buffers treated as shared memory
- Communicate + compute generate memory accesses in shared address space

Hemangee K. Kapoor

5

+34

AG



DG



SUBRATA ROY



NISHANK SIDDHARTH

NS

NISHITHA SHARMA



HARSH GUPTA

G

VA

VARHADE AMEY ANANT

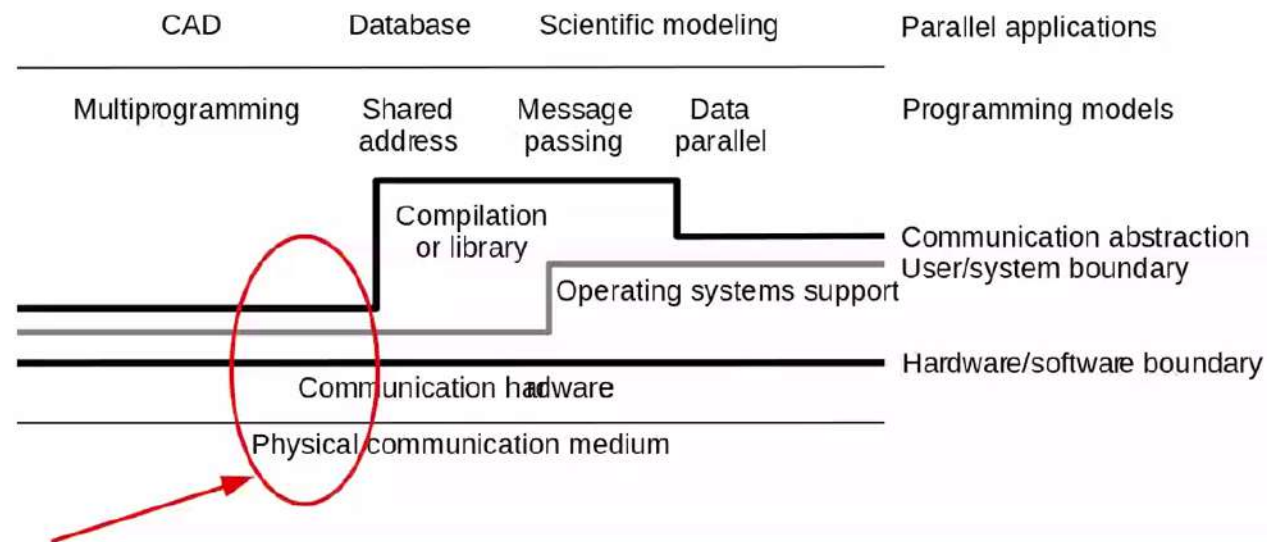
A

AMIT




Hemangee Kaipesh Kapoor


Layer Perspective




Natural extensions of memory system



Shared Cache



Centralized Memory
Dance Hall, UMA



Distributed Memory (NUMA)

