



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



How to construct the total order?

- We can construct the (hypothetical) total order that satisfies coherence by observing the following partial orders imposed by the protocol:
 - A memory operation M2 is subsequent to a memory operation M1 if the operations are issued by the same processor and M2 follows M1 in program order
 - A read operation is subsequent to a write operation W if the read generates a bus transaction that follows that for W
 - A write operation is subsequent to a read or write operation M if M generates a bus transaction and the bus transaction for the write follows that for M
 - A write operation is subsequent to a read operation M if the read does not generate a bus transaction (is a hit) and is not already separated from the write by another bus transaction

Hemangee K. Kapoor

30

+14

AS



SUBRATA ROY



NISHANK SODHARTI



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV




Hemangee Kalpesh Kapoor



How to construct the total order?

- Any serial order that preserves the resulting partial order is coherent
- The subsequent ordering relationship is transitive
- Several read transaction happening inside the processor which do not appear on the bus are not ordered by the bus
- In fact any interleaving of read operations in these internal segments is a valid serial order, as long as it obeys the program order





Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



3 states: MSI

Hemangee K. Kapoor

32

+14

AS



DG



SUBRATA ROY



NISHANK SODHARTH

NS

NISHITA SHARMA

SA

SHIVAM KUMAR AGRAWAL

SS

Syam Sarikar

SP

SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor




Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



3-state MSI write-back Invalidate Protocol



Hemangee K. Kapoor

33

+14

AS



DG



SUBRATA ROY



NISHANK SODHARTH

NS

NISHTHA SHARMA

SA

SHIVAM KUMAR AGRAWAL

SS

Syam Sarikar

SP

SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



MSI Transactions

- **Bus Read (BusRd)**
 - Generated if **PrRd** misses in the cache
 - Another cache or memory provides the block
- **Bus Read Exclusive (BusRdX)**
 - Generated by **PrWr** for block not-present or valid but read-only
 - Another cache or memory provides block (if not present in this cache)
 - Other cached copies invalidated
- **Bus Write Back (BusWB or Flush)**
 - Generated by cache controller on a **write-back**
 - The processor does not know about this
 - Main memory updated by this block
 - Write-back may be generated due to block replacement or another processor wanting to read/write block
- **Block in 'S' and incurs a write-miss**
 - **Send BusRdX**: may result in data receipt. Ignore the data as the data is already present
 - **Alternative is to send BusUpgr** transaction, where others invalidate but in response other cache/memory do not send data block on bus

Hemangee K. Kapoor

34

+14

AS



SUBRATA ROY



NISHANK SODHARTH



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



13 August 2021 - Google * 20_Aug - Mentimeter +

jamboard.google.com/d/17PzgLpIDMGmW0zBQkYwX3AP05rpt5IPyOMgXlJmKxj8/viewer?f=10

13 August 2021

11/11

Set background Clear frame

MSI

wr-Back

Inv

M = modified → dirty

S = shared → read only

I = invalid

{V, I}

+14

AS



DG



SUBRATA ROY



NISHANK SODHARTI

NS

NISHTHA SHARMA

SA

SHIVAM KUMAR AGRAWAL

SS

Syam Sarikar

SP

SARASWATULA PHANI SA/ PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



13 August 2021 - Google x 20_Aug - Mentimeter +

jamboard.google.com/d/17PzgLpIDMGmW0zBQkYwX3AP05rpzSlPyOMgXlJmKxj8/viewer?f=10

13 August 2021

11/12

Set background Clear frame

MSI

W_R-Back

Inv

M = modified → dirty

S = shared → read only

I = invalid

{V, I}

W_R-Miss

flush W_R-Back (3)

BusRd (1)

BusW_R ✓

P_RW_R

P_RRd

BusRdX (2)

+27

AS



SUBRATA ROY



DEVANSHI GUPTA



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



MSI Transactions

- **Bus Read (BusRd)**
 - Generated if **PrRd** misses in the cache
 - Another cache or memory provides the block
- **Bus Read Exclusive (BusRdX)**
 - Generated by **PrWr** for block not-present or valid but read-only
 - Another cache or memory provides block (if not present in this cache)
 - Other cached copies invalidated
- **Bus Write Back (BusWB or Flush)**
 - Generated by cache controller on a **write-back**
 - The processor does not know about this
 - Main memory updated by this block
 - Write-back may be generated due to block replacement or another processor wanting to read/write block
- **Block in 'S' and incurs a write-miss**
 - **Send BusRdX**: may result in data receipt. Ignore the data as the data is already present
 - **Alternative is to send BusUpgr transaction, where others invalidate but in response other cache/memory do not send data block on bus**

Hemangee K. Kapoor

34

+27

AS



SUBRATA ROY



DEVANSHI GUPTA



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Microsoft Teams

Search



Activity

Chat

Teams

Assignments

Calendar

Calls

...

Apps

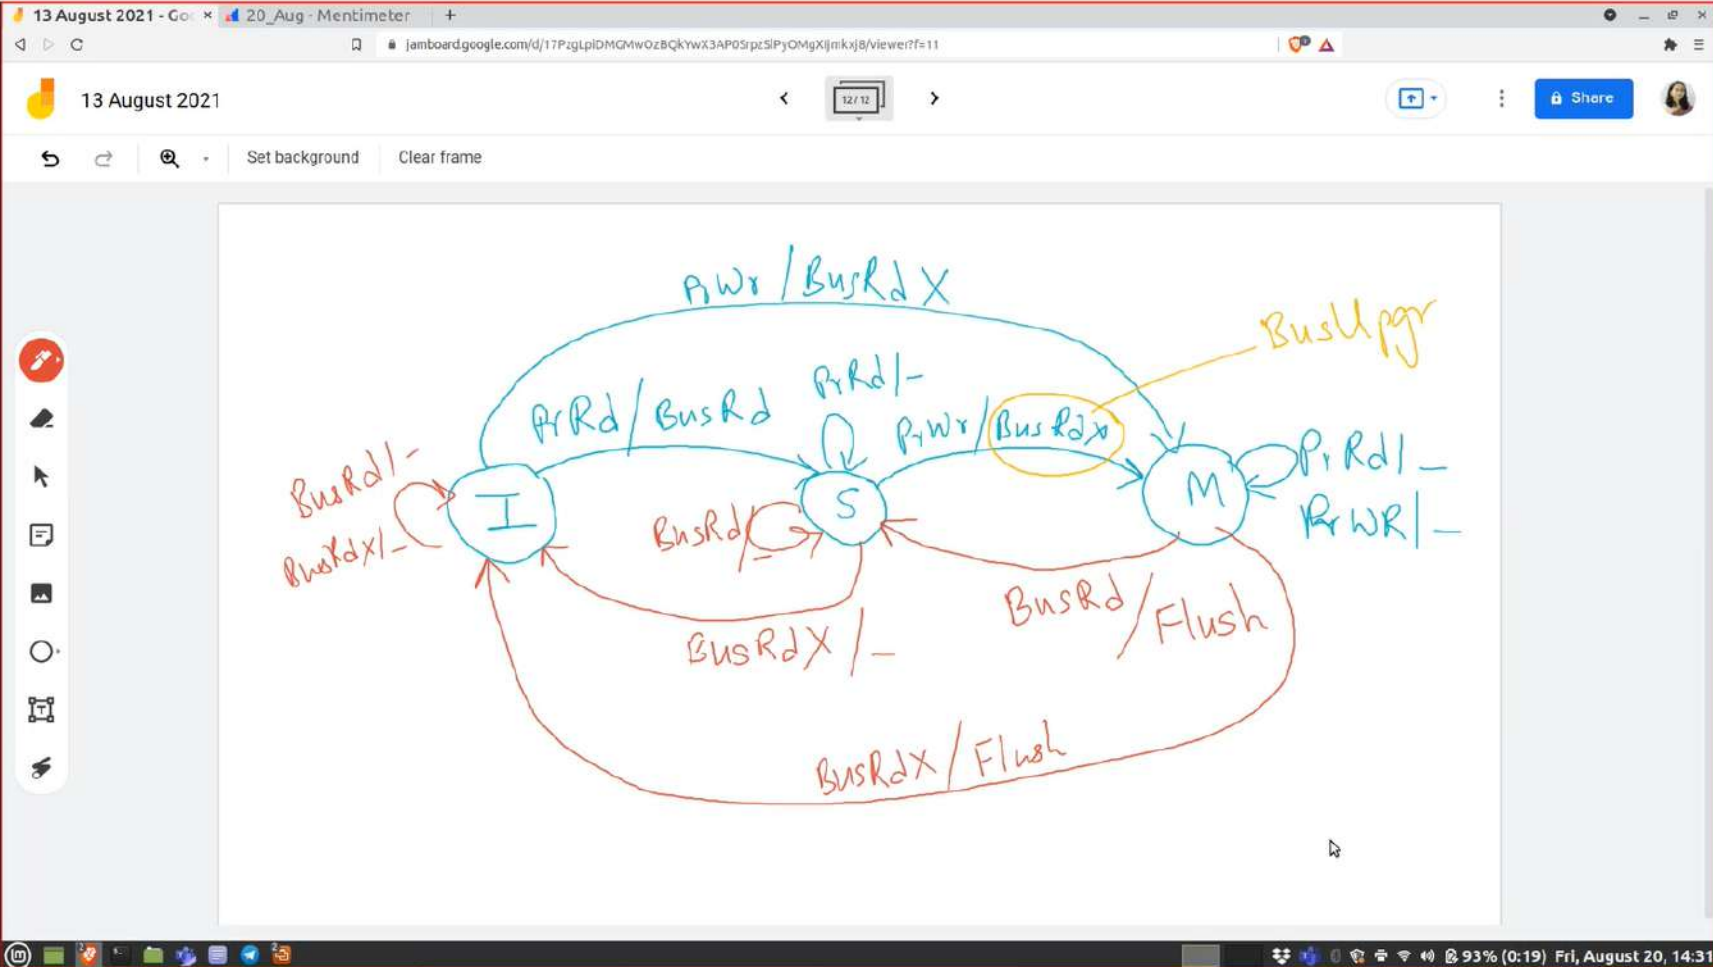
Help



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close

...



93% (0:19) Fri, August 20, 14:31

+26



YK

YOGESH KUMAR



TARVISH



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SA/ PRANAV



Hemangee Kalpesh Kapoor



Microsoft Teams

Search



Activity

Chat

Teams

Assignments

Calendar

Calls

...

Apps

Help

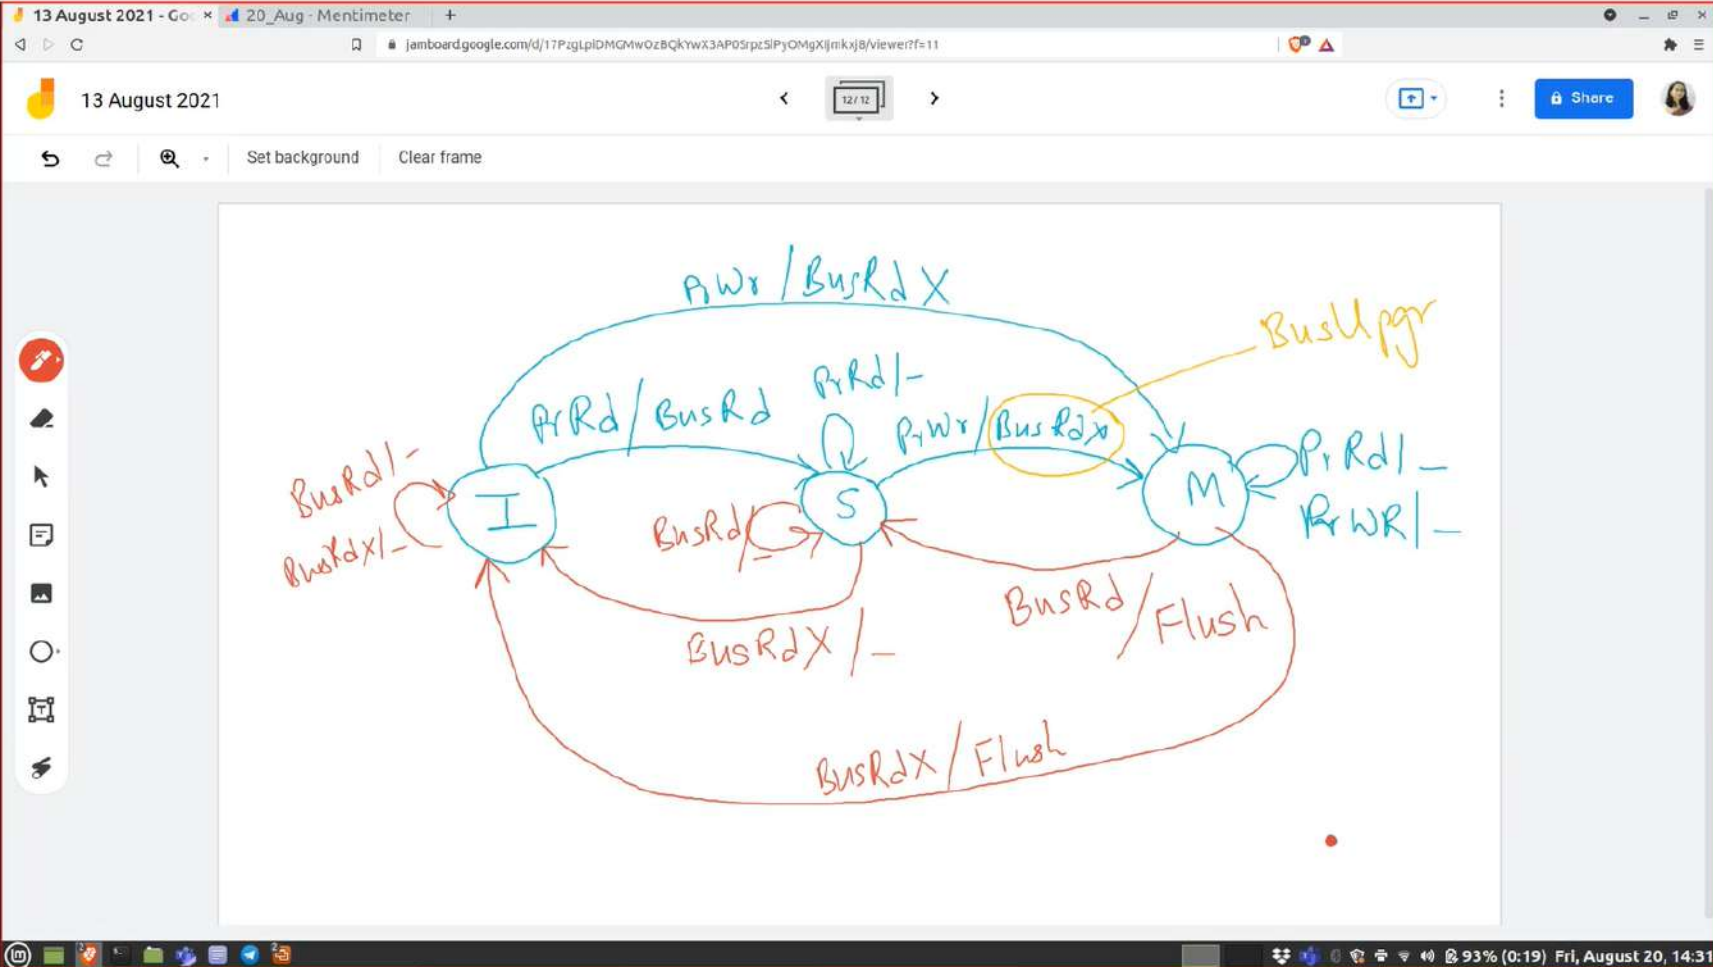
Download

Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close

...

...



93% (0:19) Fri, August 20, 14:31

+26



YK

YOGESH KUMAR



TARVISH

NS

NISHTHA SHARMA

SA

SHIVAM KUMAR AGRAWAL

SS

Syam Sarikar

SP

SARASWATULA PHANI SA/ PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



MSI state transition table

State	This Processor		Other Processor	
	LOAD	STORE	LOAD	STORE
I Invalid	Miss → S	Miss → M	-	-
S Shared	Hit	Upgrade-miss → M	-	→ I
M Modified	Hit	Hit	Send Data → S	Send Data → I

Hemangee K. Kapoor

36

+26



YK

YOGESH KUMAR



TARVISH

NS

NISHTHA SHARMA

SA

SHIVAM KUMAR AGRAWAL

SS

Syam Sarikar

SP

SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



Example: MSI

Proc-A

Proc-B

LD x

ST x

LD x

ST x

List actions taken by each processor.
Its state in local cache
Its actions taken on bus

Hemangee K. Kapoor

37

+26



YK



NS

SA

SS

SP



YOGESH KUMAR

TARVISH

NISHTHA SHARMA

SHIVAM KUMAR AGRAWAL

Syam Sarikar

SARASWATULA PHANI SAI PRANAV

Hemangee Kalpesh Kapoor




Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



Example: MSI



List actions taken by each processor.
Its state in local cache
Its actions taken on bus

- LD by Proc-B
generates BusRd,
Proc-A sends copy of A
and becomes 'S'
- ST by Proc-B
generates BusRdX,
Proc-A transitions to 'I'





Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



MSI protocol satisfies Coherence and Consistency

- Satisfies Consistency will be done towards the end ... later in the course
- Here we will prove Coherence
- Satisfying Coherence requires to prove
 - (1) Write Propagation
 - (2) Write Serialisation

Hemangee K. Kapoor

38

+26



YOGESH KUMAR



TARVISH



NISHITA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



Write Propagation correctness

- Cached block is updated several times in wr-back protocol before memory is actually updated
- Read by another processor obtains block from this cache and not memory. The "flush" will also update memory
- Write-hits NOT visible on bus
- **When is write-miss made visible to others?**
 - Write to 'S' or 'I' block causes BusRdX transaction
 - Writer observes that data in its cache after this transaction
 - Write will be made visible to other processors by Invalidation due to BusRdX, and those processors will experience a cache miss before actually observing the written value
- **When is write-hit visible?**
 - Write hits can be observed by other processors only after a cache miss through a bus transaction

Hemangee K. Kapoor

39

+26



YOGESH KUMAR



TANVISH



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



Write Serialisation correctness

- **Writes that appear on bus (BusRdX) are ordered by Bus**

- BusRdX guarantees that only one valid copy is present to write. No one read it anymore.
- Write performed in writer's cache before other transactions => so ordered in same way for all processors (including the writer) wrt other bus transaction
- Multiple writes can happen to same location by writer, however the key difference is that between two transactions for that block that appear on the bus, only 1 processor can perform such write hits
- And this processor (P) is the one who did the write. **For P, multiple writes appear in program order.** In serialisation (or serial order), these write-hits appear in program order before the next bus transaction.

- **Writes serialised wrt Reads**

- Processor P issues BusRdX
- Further, memory operations are done by P (read/write hit) in program order
- Therefore reads by P will see them in this (program) order wrt other writes (made by P)
- For read by another processor (P2) there is at least one bus transaction that separates the completion of this read (by P2) from the completion of the write hits (by P)
- Therefore bus transaction ensures that, the read (by P2) also sees the write in the same consistent serial order

=> Reads by all processors see all writes in the same order

Hemangee K. Kapoor

40

+26



YOGESH KUMAR



TANVISH



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SA/ PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



4 states: MESI

Hemangee K. Kapoor

41

+26



YOGESH KUMAR



TARVISH



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



4-state MESI write-back, Invalidate Protocol

- Concern in MSI protocol:
 - If a sequential application is running on multiprocessor, i.e. no sharing of data
 - We still need 2 bus transactions in MSI for Rd / Wr
 - 1st a BusRd to load block from memory in 'S' state
 - 2nd a BusRdX (or BusUpgr) that converts 'S' -> 'M'
- Therefore add another state 'E'
 - Exclusive clean = Exclusive un-owned
 - Memory also has updated copy
- Exclusive = writable = No sharers, so write without informing others
- Modified = owned = written

Hemangee K. Kapoor

42

+26



YOGESH KUMAR



TANVISH



NISHTHA SHARMA



SHIVAM KUMAR AGRAWAL



Syam Sarikar



SARASWATULA PHANI SAI PRANAV



Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



4-state: MESI

- States
 - Invalid
 - Exclusive: Only this cache has a copy and not modified
 - Shared: Two or more caches may have copy
 - Modified: Dirty
- FSM shown on another slide
- Variants of this used in Intel Pentium, PowerPC, MIPS 4400 of Silicon Graphics Challenge multiprocessor
- It was first published by researchers at UIUC and hence it is often called the Illinois protocol

Hemangee K. Kapoor

43

+23

AS



DG

YK



NS



SS

SP



YOGESH KUMAR

TARVISH

NISHTHA SHARMA

SUBRATA ROY

Syam Sarikar

SARASWATULA PHANI SAI PRANAV

Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



I -> E ??

- On PrRd we transition from 'I' to 'E' if no cache has a copy of the block
- How can you tell? => Need hardware support
- All cache controllers snoop on BusRd
- Assert additional signal: 'shared' signal if cache has the block. This is the variable 'S' and 'S-bar' in the FSM
- Wired-OR signal, therefore any cache making 'shared'=1 we know that we should perform I->S
- And if 'shared'=0 then we should perform I->E

Hemangee K. Kapoor

44

+23

AS



DG

YK



NS



SS

SP



YOGESH KUMAR

TANVISH

NISHTHA SHARMA

SUBRATA ROY

Syam Sarikar

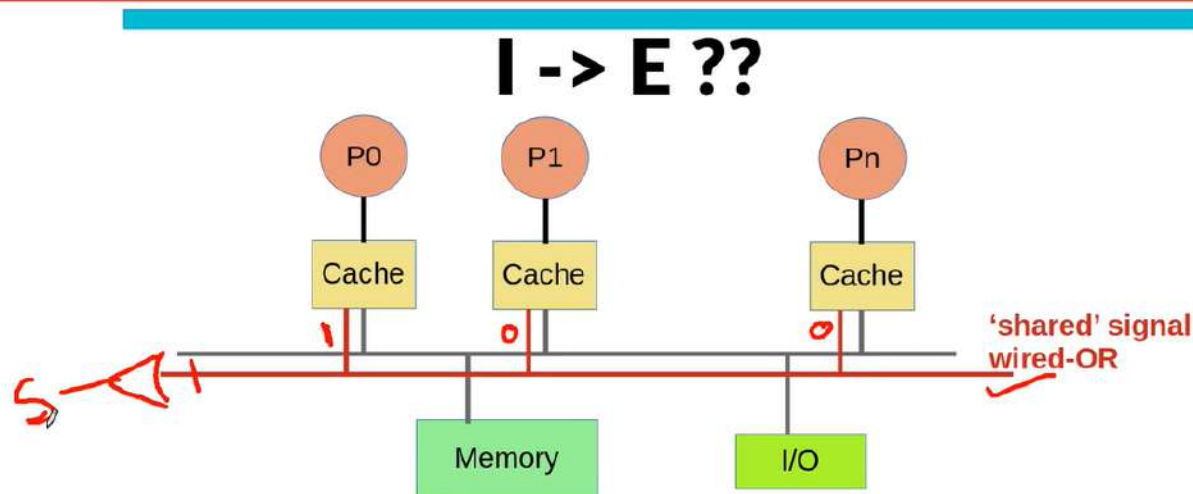
SARASWATULA PHANI SAI PRANAV

Hemangee Kalpesh Kapoor



Meeting in _General_-20210820_140419-Meeting Recording.mp4

Close



- BusRd(**S**) means shared line is asserted on BusRd transaction
- BusRd(**S-bar**) means 'S' is un-asserted
- **Flush** : if cache-to-cache transfers: only one cache flushes data, giving it to requestor
- **Satisfying Coherence**: Arguments same as MSI
- **MOESI** protocol: Owned state = Exclusive but memory not valid
 - M->S : But one cache has updated copy
 - Memory not updated. Ownership transferred. Cache in state 'O' provides the data
- It is possible to have a block in state 'S' even when there are not other sharers
 - Why? Because, others removed their copies and deletion does not trigger transaction. This cache may not know it is the only one left with a copy

Hemangee K. Kapoor

45

+23

AS



DG

YK



NS



SS

SP



YOGESH KUMAR

TANVISH

NISHTHA SHARMA

SUBRATA ROY

Syam Sarikar

SARASWATULA PHANI SAI PRANAV

Hemangee Kalpesh Kapoor