

# Why are we building parallel systems?

- Single processor performance
  - Transistor density
  - IC
  - Speed
  - ILP
- Speed increases, power increases, heat increases
- Chips too hot to be reliable. Limits to heat dissipation
- How to exploit the increasing number of transistor density?
  - Add parallelism
  - Go multicore




# Why we need to write parallel programs?

- Programs written for conventional single core systems cannot exploit presence of multiple cores
- We can run multiple instances of a program on a multicore system. BUT this is of less help
- e.g. running multiple instances of gaming program? Is not desirable !
  - We want the game to run faster !
- Rewrite serial programs to parallel programs so that they use multiple cores
- Rewrite or translate serial code
  - Translate => can we auto-translate?
  - Till now researchers have had limited success in writing programs that convert the serial program to parallel program
  - It is not straightforward to convert serial to parallel by identifying par-constructs
- Rather, step back and devise entirely new algorithm



# Example: sum of array elements



Hemangee K. Kapoor

13

# Example

- Sum of all elements of an array
- Tough for the translator program to discover the procedure
- Certain softwares can identify common serial constructs and efficiently parallelise them
- HOWEVER to apply the principles on ever increasing complex serial programs is difficult
- Therefore, we cannot simply continue to write serial programs.
- We MUST write parallel programs to exploit the power of multiple cores



# How do we write parallel programs

- Basic idea is of partitioning the work to be done among the cores
  - Task parallelism
  - Data parallelism
- Ex: Prof. P, TAs: A, B, C, D
  - Have to check 100 answer books with 5 questions each
  - Data parallel  $\Rightarrow$  P, A, B, C, D each take 20 copies to check
  - Task parallel  $\Rightarrow$  P, A, B, C, D each take 1 question each to check



# In theory

- Sequential
  - Time to sum  $n$  numbers ?
  - $O(n)$
  - Time to sort  $n$  numbers ?
  - $O(n \log n)$
- Parallel
  - Time to sum  $n$  numbers?
  - Tree for  $O(\log n)$
  - Time to sort  $n$  numbers
  - Non trivially  $O(\log n)$





# But in practice

- Name a datum across processors?
- Communicate values?
- Coordinate and synchronize?
- Select processing node size (few-bit ALU to a PC)?
- Select number of nodes in system?

# Programming Model

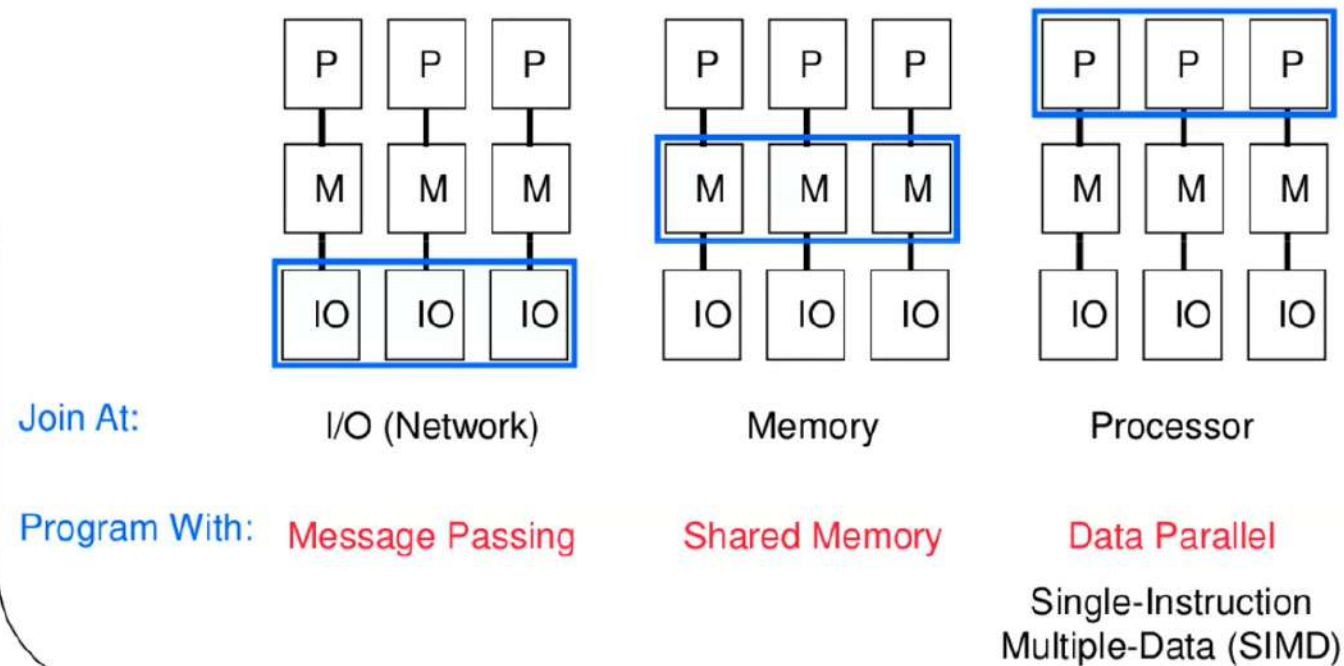
- Provides a communication abstraction that is a contract between hardware and software (a la ISA)
  - Programming model != programming language
- Conceptualisation of the machine that the programmer uses in coding applications
  - How parts cooperate and coordinate their activities
  - Specifies communication and synchronisation operations
- **Multiprogramming**
  - No communication or synchronisation at program level
- **Shared address space**
  - Like a bulletin board
- **Message Passing**
  - Like letters or phone calls, explicit point of contact
- **Data Parallel**
  - More regimented, global actions on data
  - Implemented using either shared address space or message passing





## Historical View


- Historically: system architecture and programming model were tied together



# Historical view ...

- Architecture -> Programming Model
  - Join at network -> program with message passing model
  - Join at memory -> program with shared memory model
  - Join at processor -> program with SIMD or data parallel
- Programming Model -> Architecture
  - Message-passing programs on message-passing arch
  - Shared-memory programs on shared-memory arch
  - SIMD/data-parallel programs on SIMD/data-parallel arch
- But
  - Isn't hardware basically the same? Processors, memory, & I/O?
  - **Convergence!** Why not have generic parallel machine & program with model that fits the problem?

## A Generic Parallel Machine



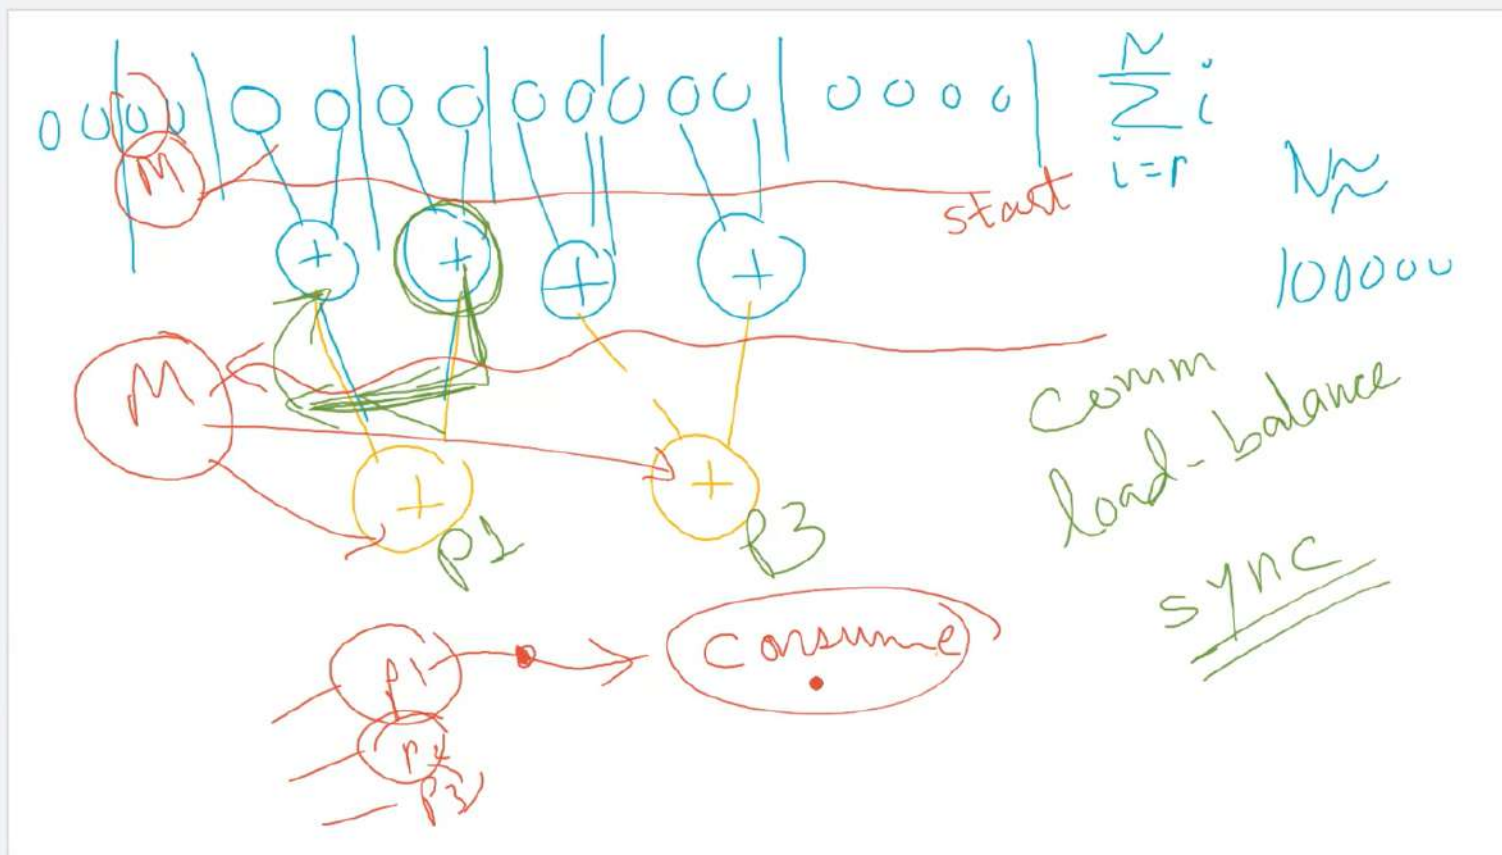
- Separation of programming models from architectures
- All models require communication
- Node with processor(s), memory, **communication assist**

Reminder: could be on one chip or many

# Coordination

- When the cores can work independently, writing a parallel program is much same as writing a serial program
- It gets complex when the cores need to coordinate their work
- Coordination involves
  - Communication
    - Send partial sums to another core
  - Load Balancing
    - We want each core to do same amount of work. Otherwise some cores are working and others are waiting idle wasting computational power
  - Synchronisation
    - Master reads the values in the array
    - Then cores must start computing. Else wait
    - Add a point of synchronisation between read value and compute partial sums







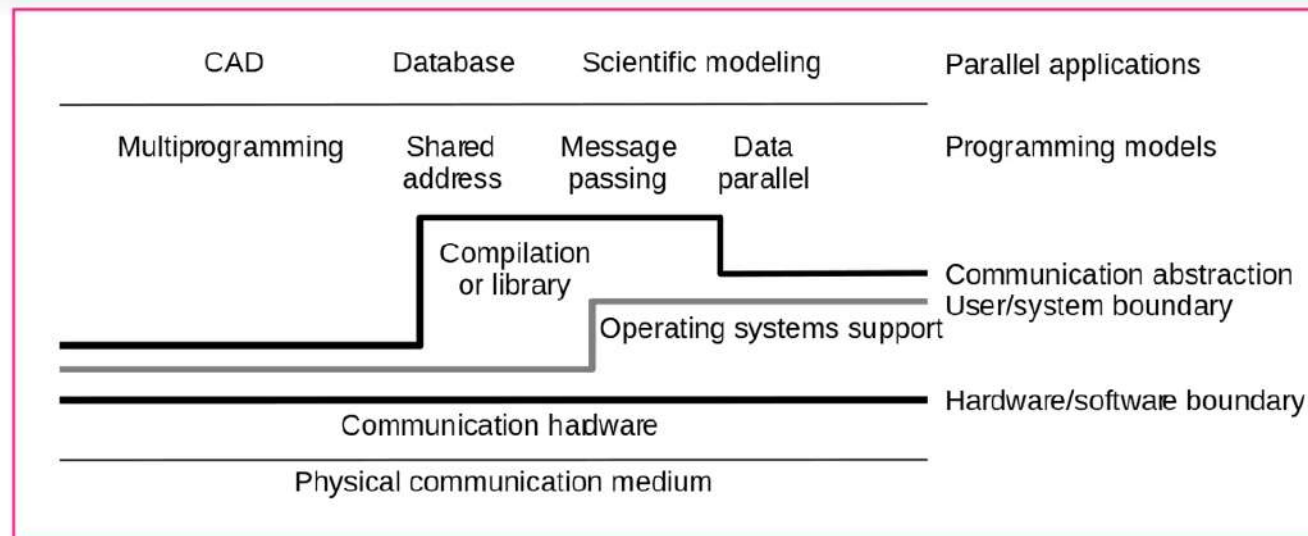
# Parallel Programming

- For Distributed Memory and Shared Memory Systems
- Using C and its extensions
- MPI: Message Passing Interface
  - For Distributed memory systems
  - Libraries of type definitions + functions + macros
- POSIX threads: pthreads
  - For shared memory systems
  - Libraries of type definitions + functions + macros
- OpenMP
  - For shared memory systems
  - Library + modifications to C compiler

# Parallel Architectures

- Almasi and Gottlieb 1989
  - “Parallel computer is a collection of processing elements that communicate and cooperate to solve large problems fast”

# Layers of Abstraction



As programming models have become better understood and Implementation techniques have matured, compilers and run-time libraries have grown to provide an important bridge between the programming model and the underlying hardware

# Layers of abstraction in parallel computer architecture

- The framework to understand communication in parallel machines is shown in the figure (another slide)
- Top layer = programming model
  - Specifies how the programs running in parallel communicate information to one another, and
  - What synchronisation operations are available to coordinate activities
- Communication abstraction = user level communication primitives
  - Provided directly by hardware, or
  - By operating system, or
  - Machine specific user software



# Layers of abstraction in parallel computer architecture

- The framework to understand communication in parallel machines is shown in the figure (another slide)
- Top layer = programming model
  - Specifies how the programs running in parallel communicate information to one another, and
  - What synchronisation operations are available to coordinate activities
- Communication abstraction = user level communication primitives
  - Provided directly by hardware, or
  - By operating system, or
  - Machine specific user software
- Hardware/Software Boundary = almost flat as most hardware features are uniform in complexity





# Shared Address Space, Interconnects, Message Passing

Hemangee K. Kapoor

9

+52

VP

VA

UK

AM



IMLIJUNGLA LONGCHAR

VK

VEDIKA JITENDRA KULKARNI



RATHOD SAINATH



BHASKER GOEL

VB

VANSHITA BANSAL

SB

SWARNIL BANDE



Hemangee Kaipesh Kapoor